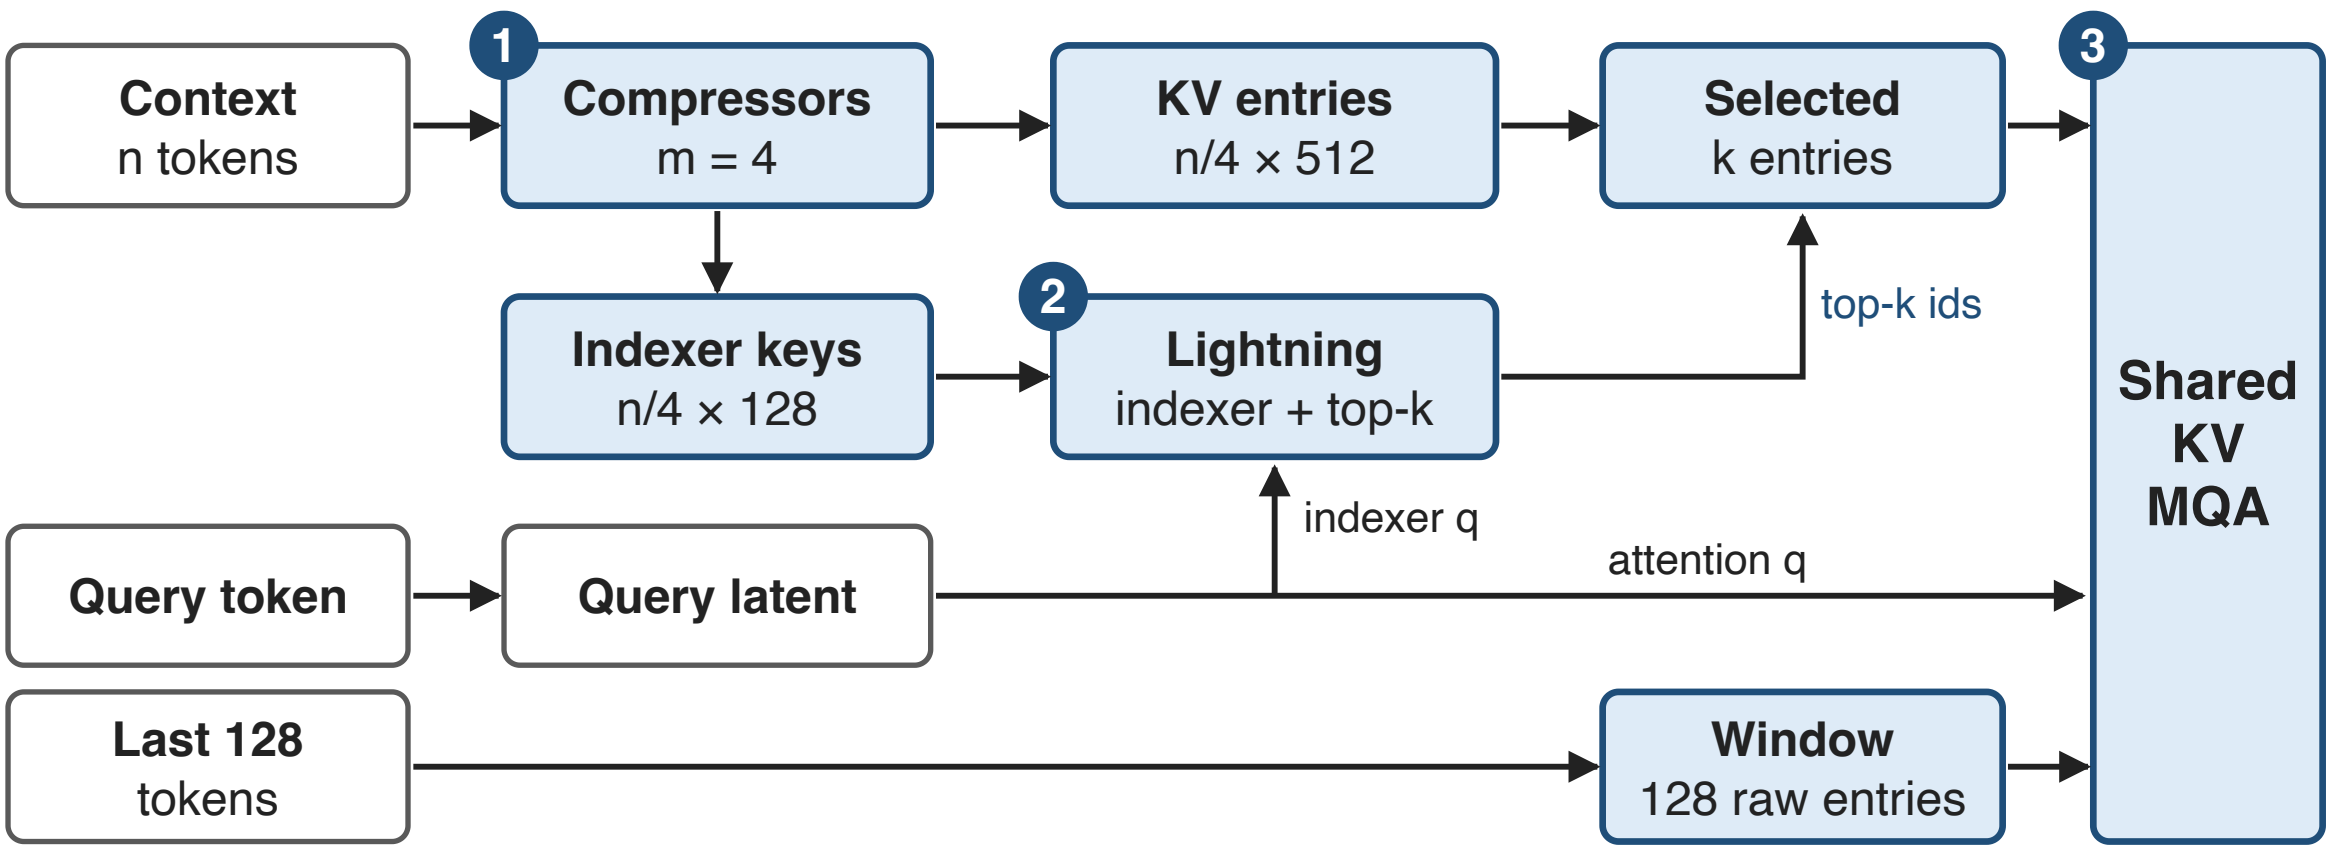
Context
n tokens
1
Compressors
m = 4
KV entries
n/4 × 512
Selected
k entries
3
Shared
KV
MQA
Indexer keys
n/4 × 128
2
Lightning
indexer + top-k
top-k ids
indexer q
attention q
Query token
Query latent
Last 128
tokens
Window
128 raw entries