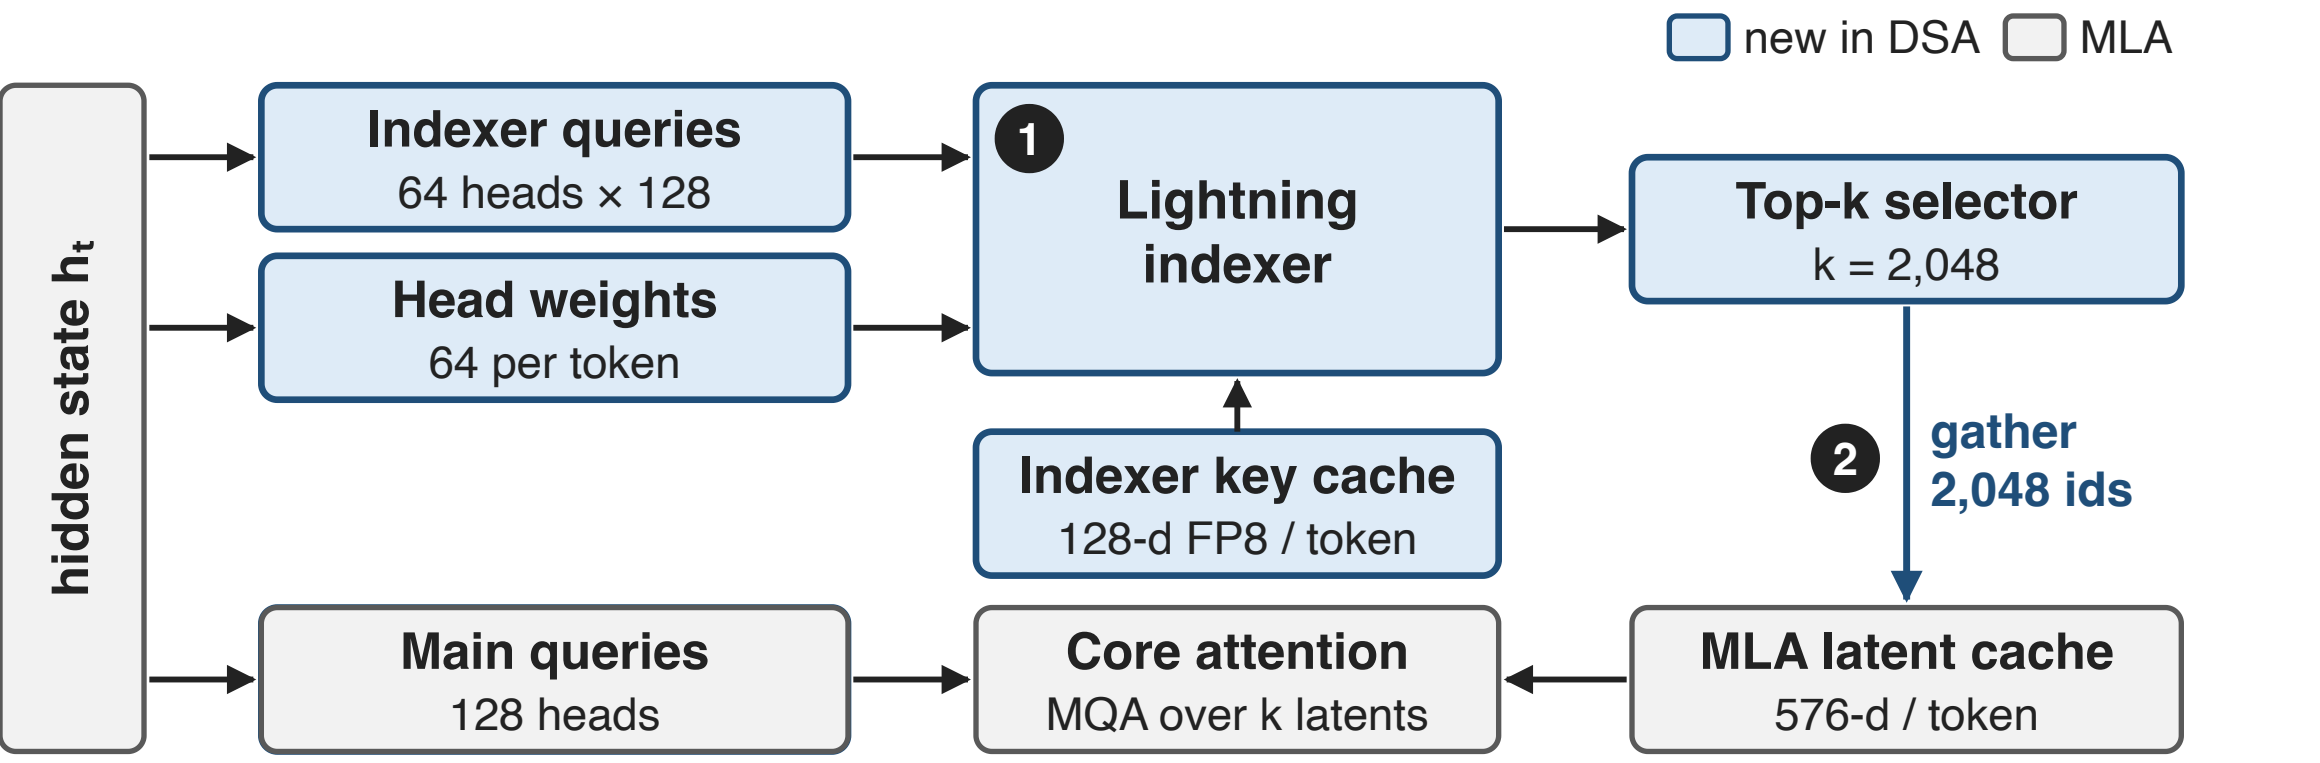
new in DSA
MLA
hidden state h_t
Indexer queries
64 heads × 128
Head weights
64 per token
1
Lightning indexer
Indexer key cache
128-d FP8 / token
Top-k selector
k = 2,048
2
gather 2,048 ids
Main queries
128 heads
Core attention
MQA over k latents
MLA latent cache
576-d / token