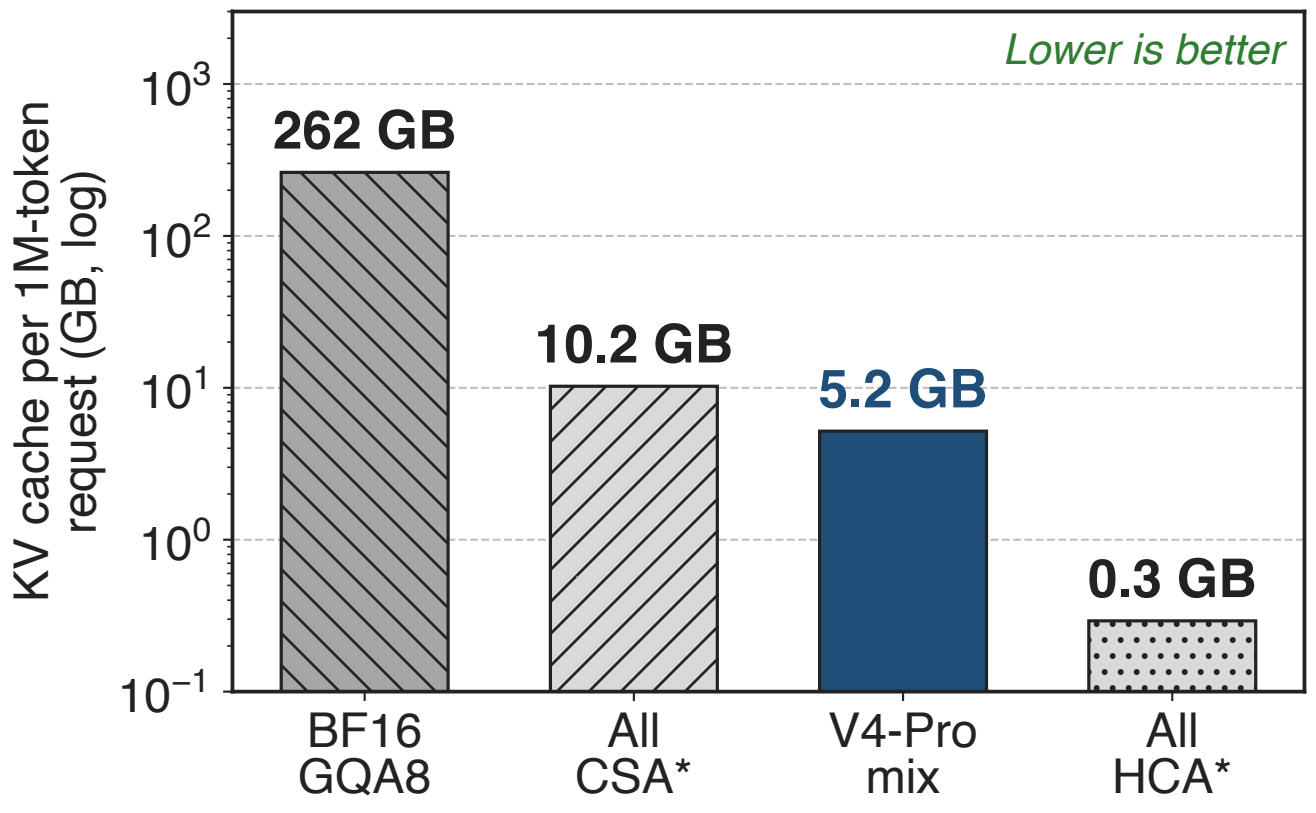
Lower is better
KV cache per 1M-token request (GB, log)
10^3
10^2
10^1
10^0
10^{-1}
262 GB
10.2 GB
5.2 GB
0.3 GB
BF16 GQA8
All CSA*
V4-Pro mix
All HCA*