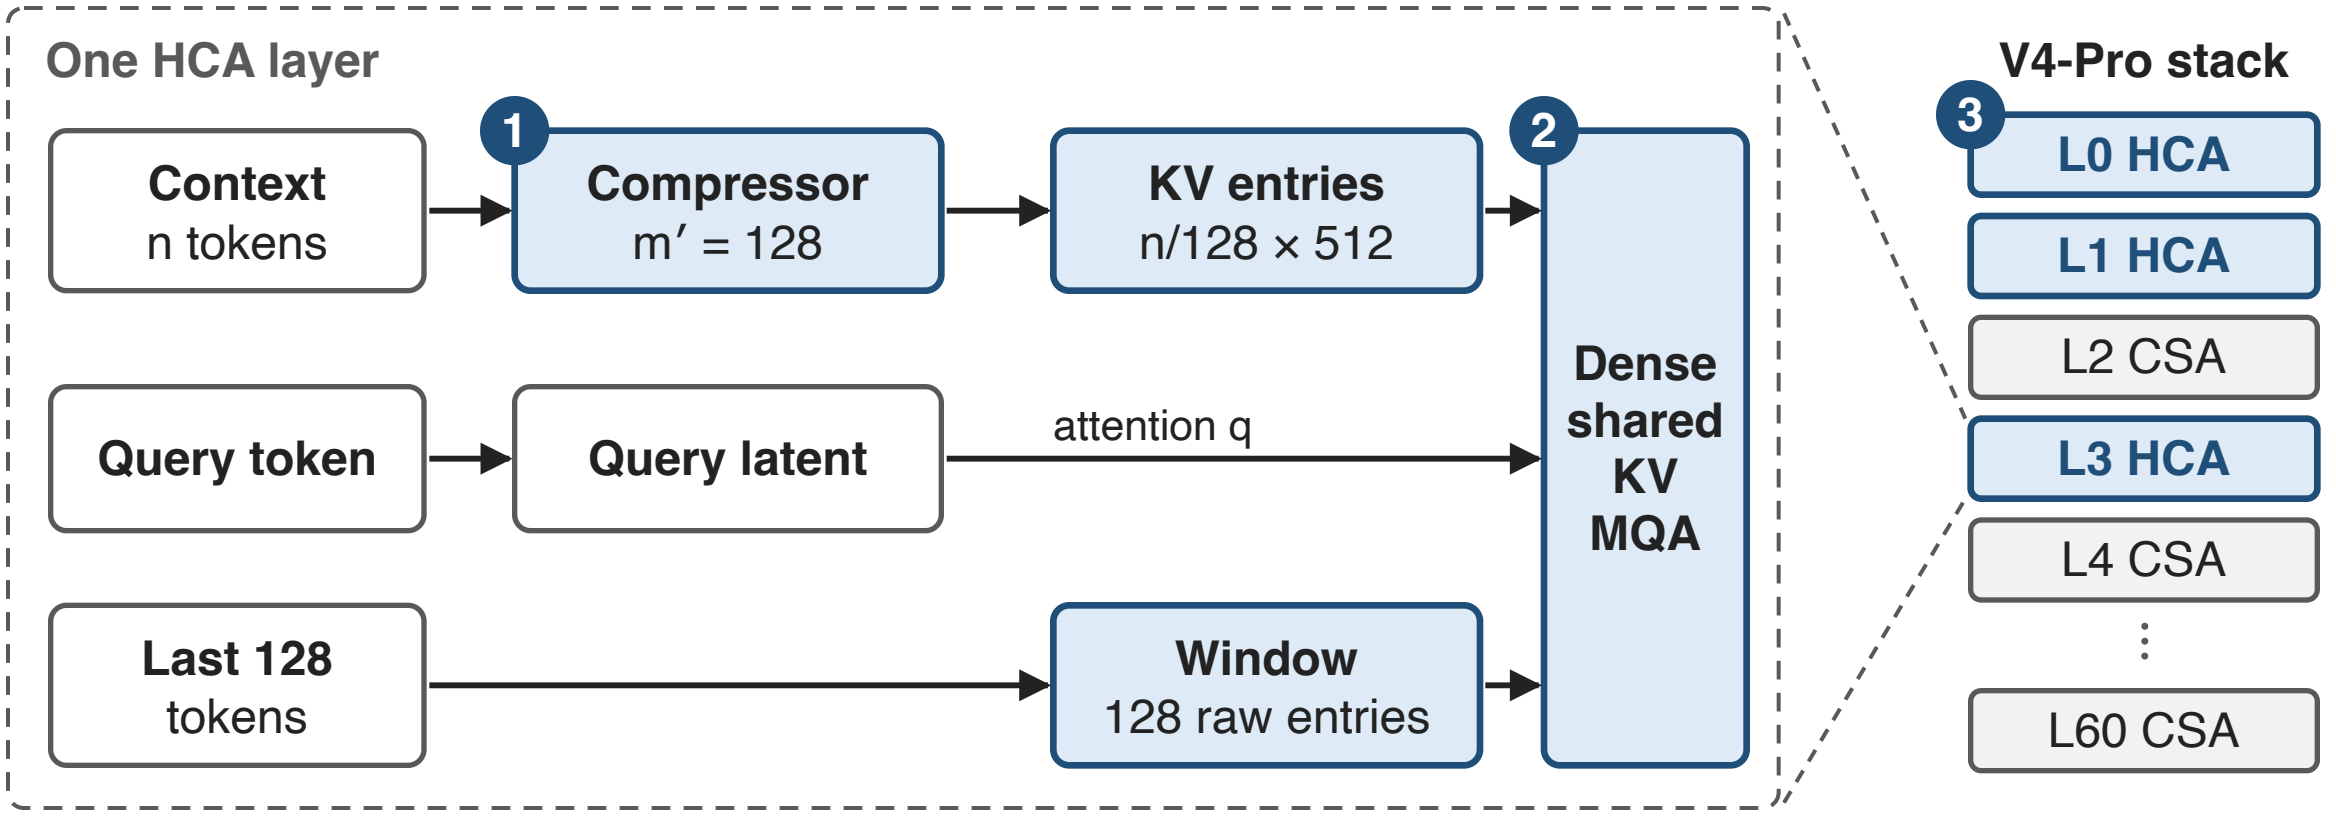
One HCA layer
Context
n tokens
1
Compressor
m′ = 128
KV entries
n/128 × 512
2
Dense
shared
KV
MQA
Query token
Query latent
attention q
Last 128
tokens
Window
128 raw entries
V4-Pro stack
3
L0 HCA
L1 HCA
L2 CSA
L3 HCA
L4 CSA
L60 CSA