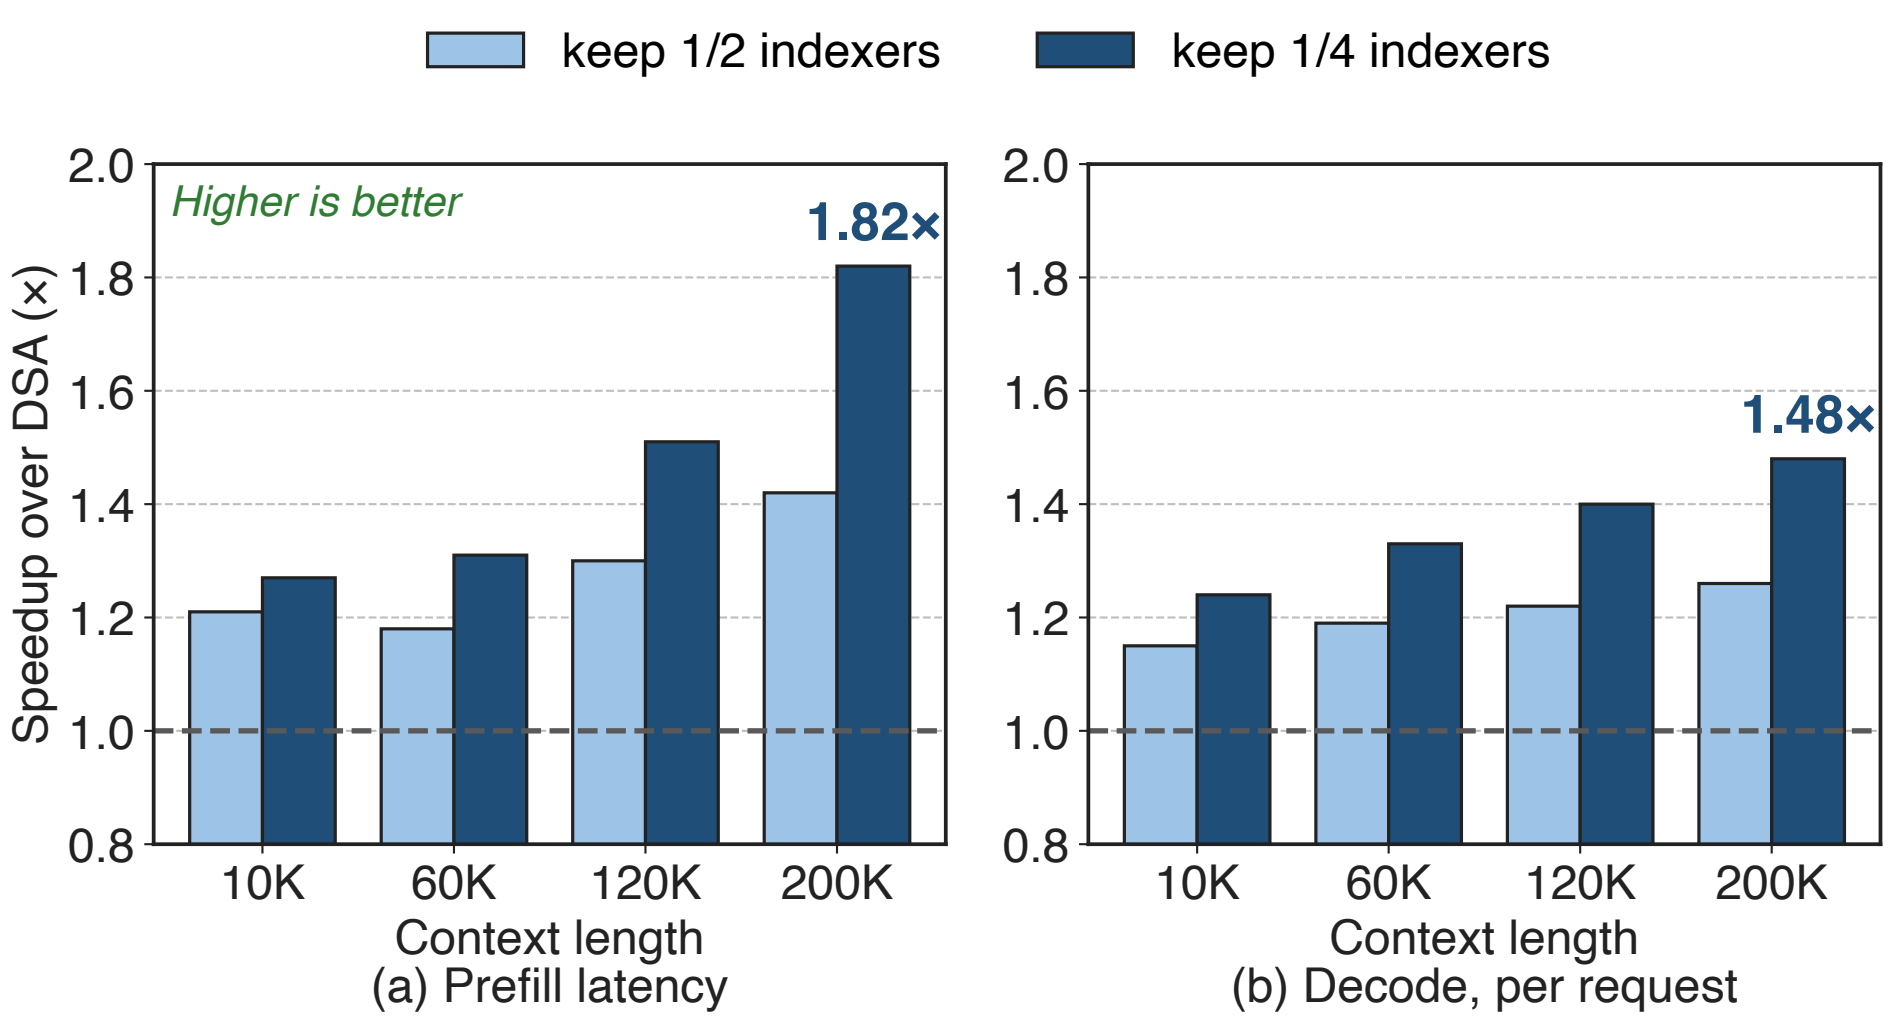
keep 1/2 indexers
keep 1/4 indexers
Higher is better
1.82×
1.48×
Speedup over DSA (x)
2.0
1.8
1.6
1.4
1.2
1.0
0.8
10K
60K
120K
200K
Context length
(a) Prefill latency
Context length
(b) Decode, per request