## In FlashKDA v1

- ✓ **Hopper and Blackwell**
  SM90+ · CUDA 12.9+
- ✓ **Head dims K = V = 128**
  only this shape
- ✓ **Bounded gate**
  lower bound in [–5, 0]
- ✓ **Forward pass and varlen**
  prefill · cu_seqlens batches

## Not in the public v1

- ✗ **Backward pass**
  K3 reports training use
- ✗ **Decoding**
  separate fused kernel (K3 §5.4.2)
- ✗ **Other head dims**
  and pre-Hopper GPUs
- ✗ **Unbounded gates**
  the API requires a lower bound