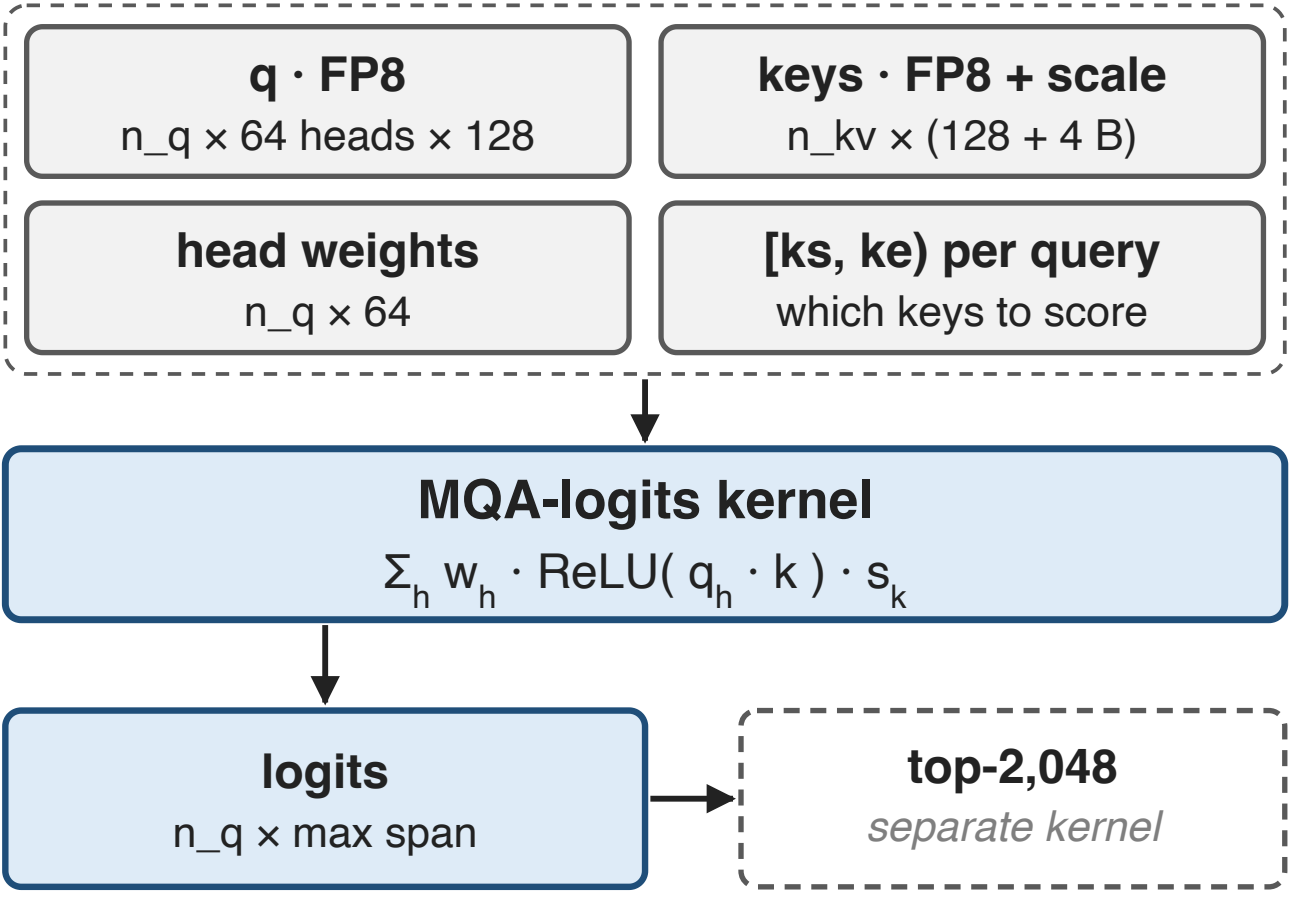
q · FP8
n_q × 64 heads × 128
keys · FP8 + scale
n_kv × (128 + 4 B)
head weights
n_q × 64
[ks, ke) per query
which keys to score
MQA-logits kernel
$\Sigma_h w_h \cdot \mathrm{ReLU}(q_h \cdot k) \cdot s_k$
logits
n_q × max span
top-2,048
separate kernel