① Layer

Split the collective

AllReduce → RS + AG

Norm and router run on sharded tokens

② Kernel

Fuse it into the matmul

Every tile ships its partial inside one kernel

③ Cluster

Split the roles

Prefill nodes → Decode nodes

Prefill cannot preempt a running decode

Same goal at three scales: no rank waiting on the fabric