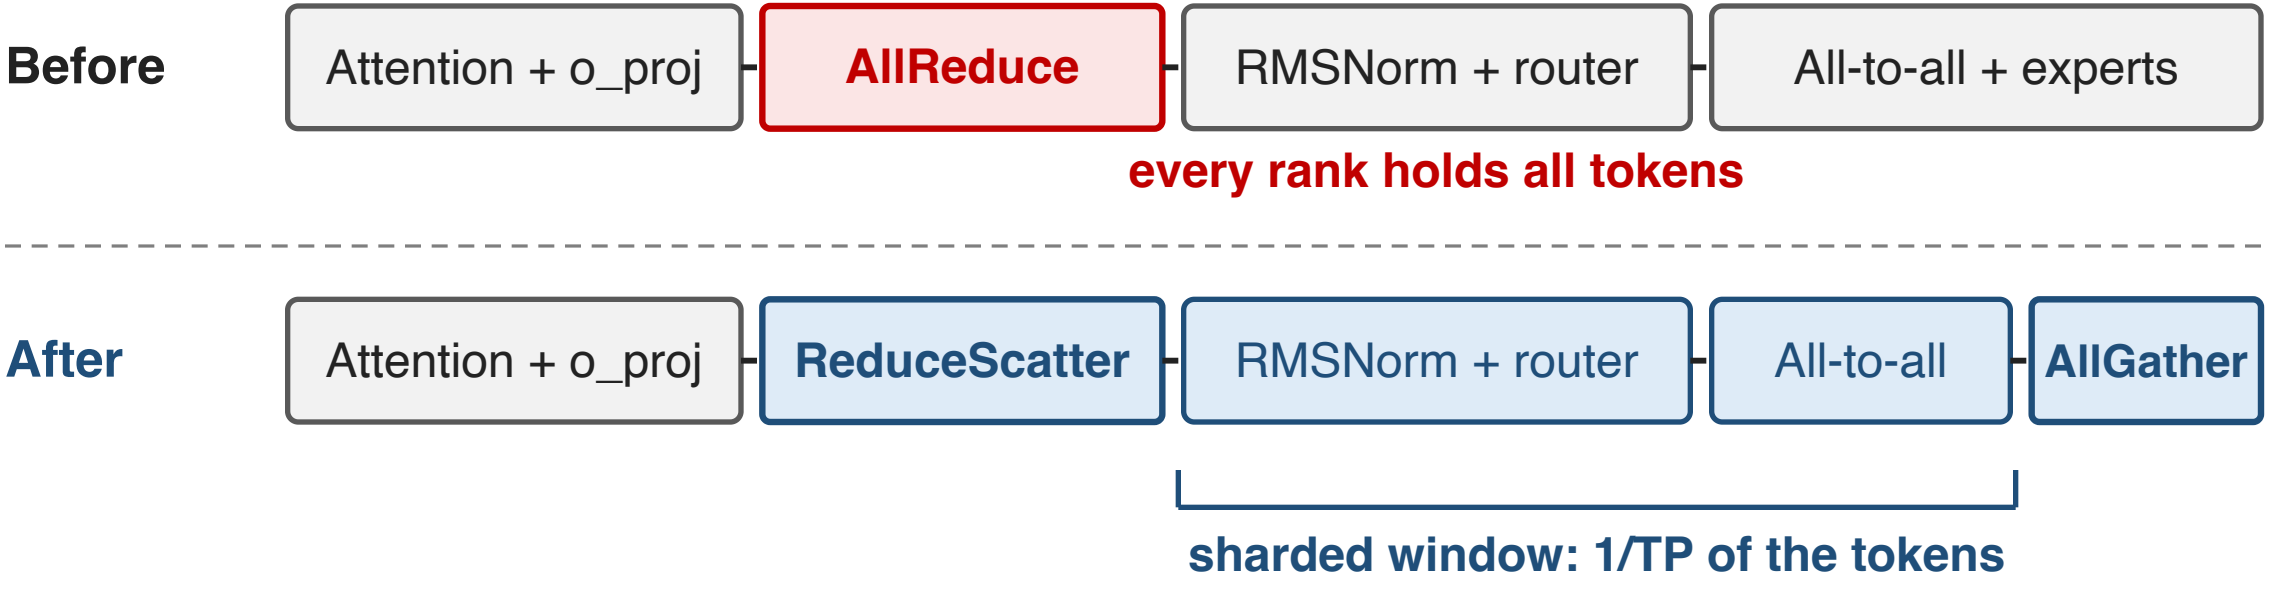
Before
Attention + o_proj
AllReduce
RMSNorm + router
All-to-all + experts
every rank holds all tokens
After
Attention + o_proj
ReduceScatter
RMSNorm + router
All-to-all
AllGather
sharded window: 1/TP of the tokens