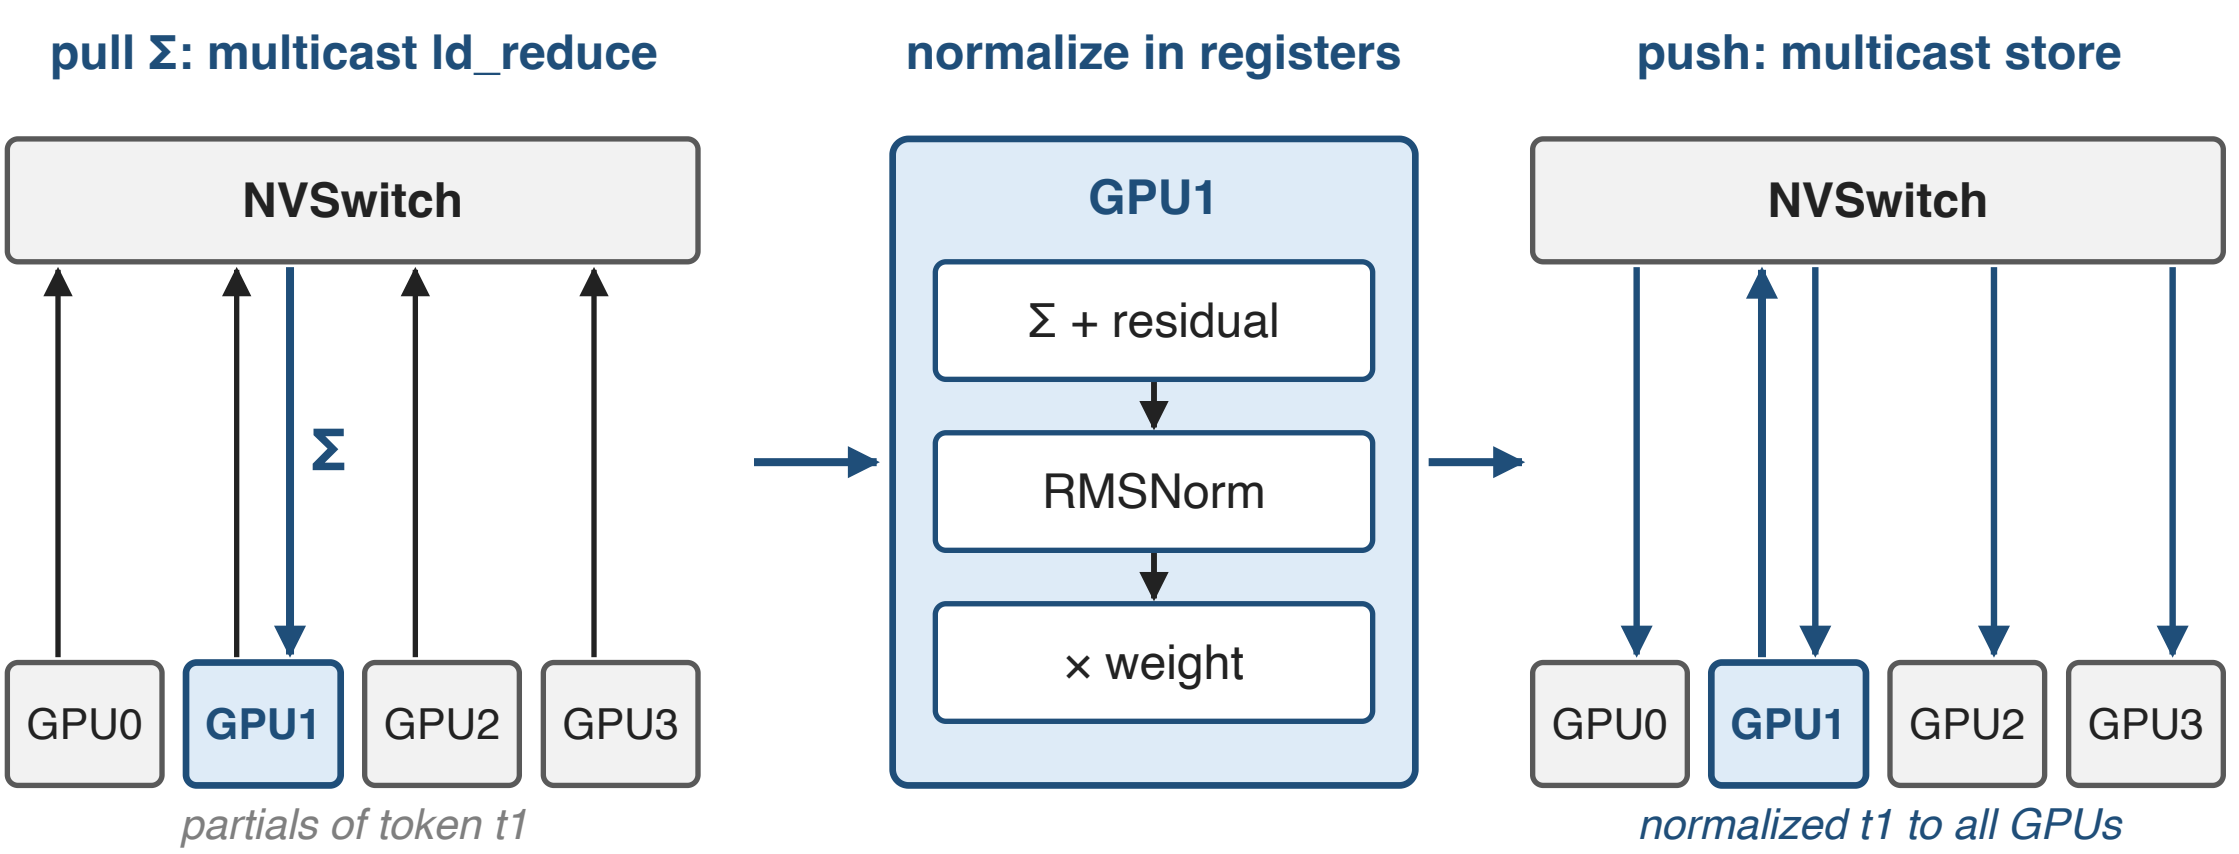
pull Σ: multicast ld_reduce
normalize in registers
push: multicast store
NVSwitch
Σ
GPU0
GPU1
GPU2
GPU3
partials of token t1
GPU1
Σ + residual
RMSNorm
× weight
NVSwitch
GPU0
GPU1
GPU2
GPU3
normalized t1 to all GPUs