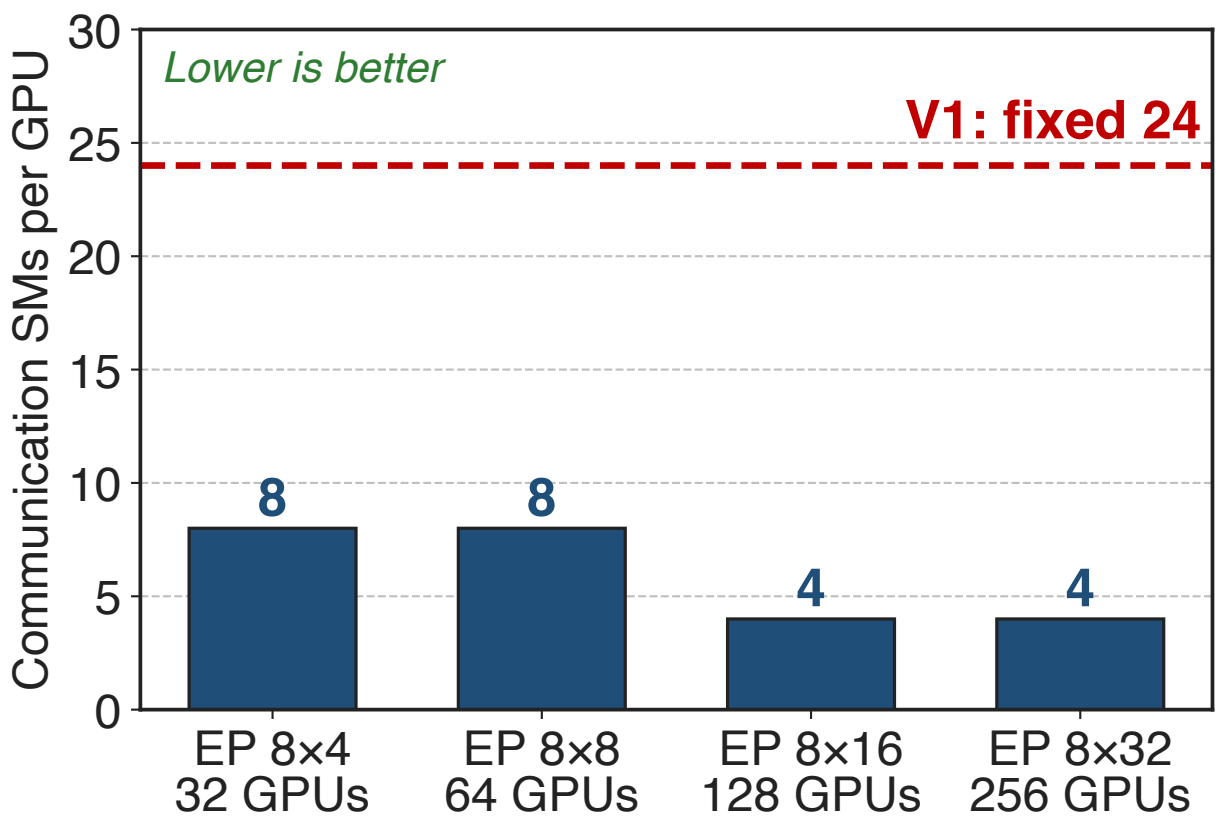
Lower is better
V1: fixed 24
Communication SMs per GPU
30
25
20
15
10
5
0
8
8
4
4
EP 8×4
32 GPUs
EP 8×8
64 GPUs
EP 8×16
128 GPUs
EP 8×32
256 GPUs