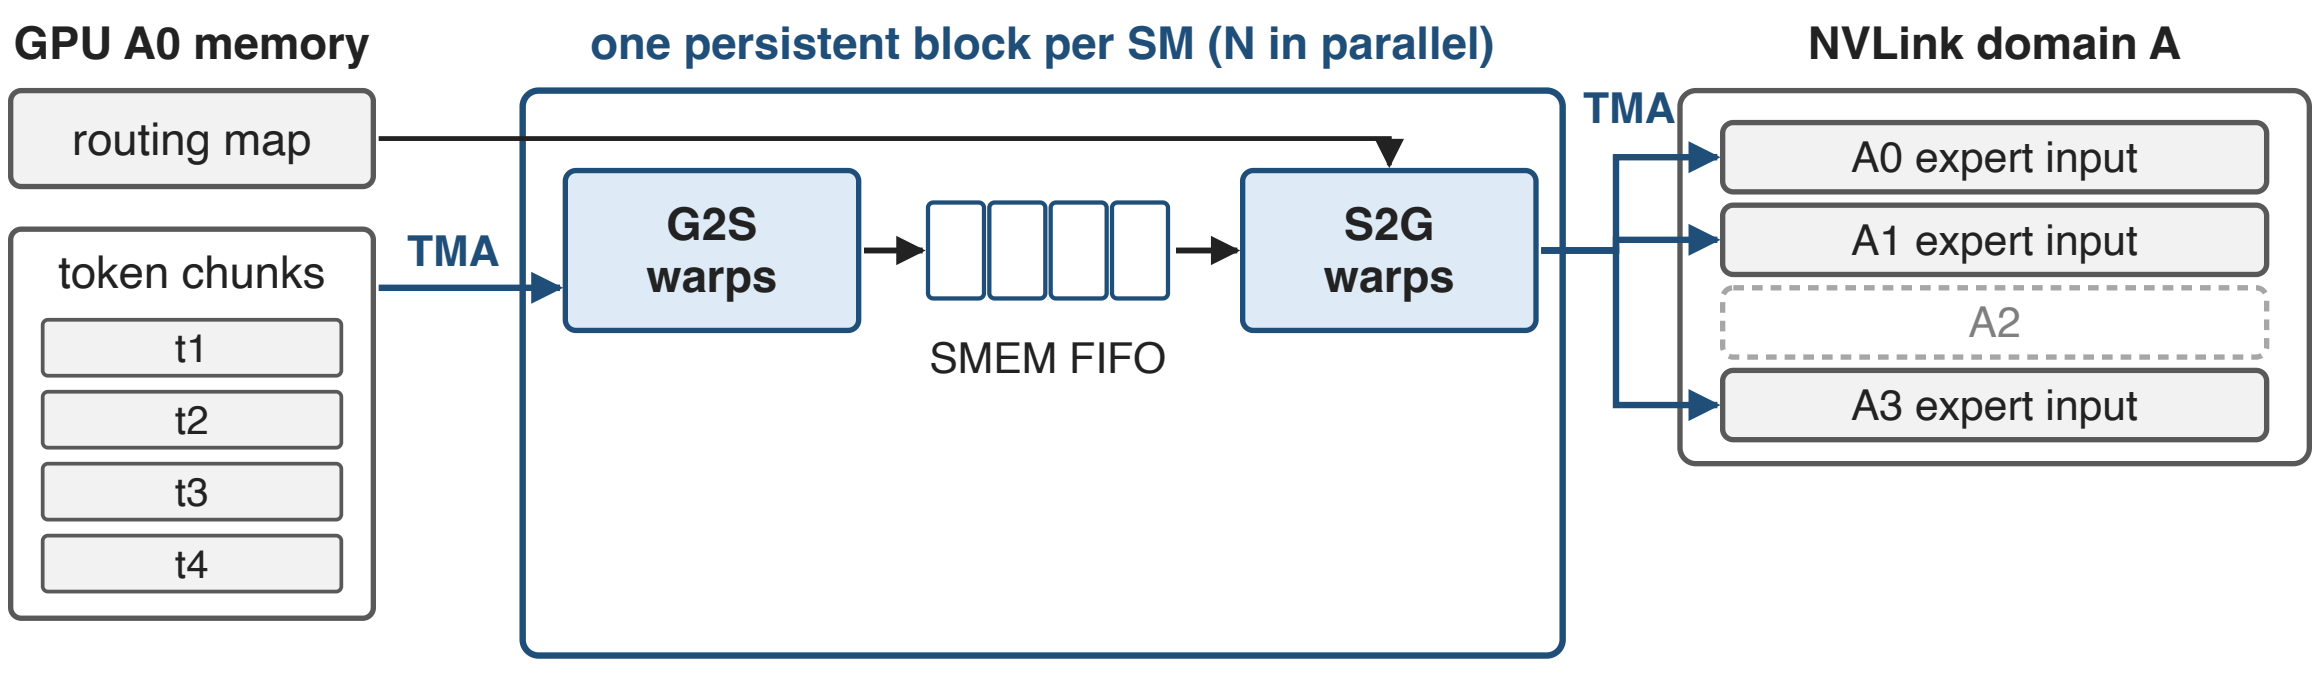
GPU A0 memory
one persistent block per SM (N in parallel)
NVLink domain A
routing map
token chunks
t1
t2
t3
t4
TMA
G2S warps
SMEM FIFO
S2G warps
TMA
A0 expert input
A1 expert input
A2
A3 expert input