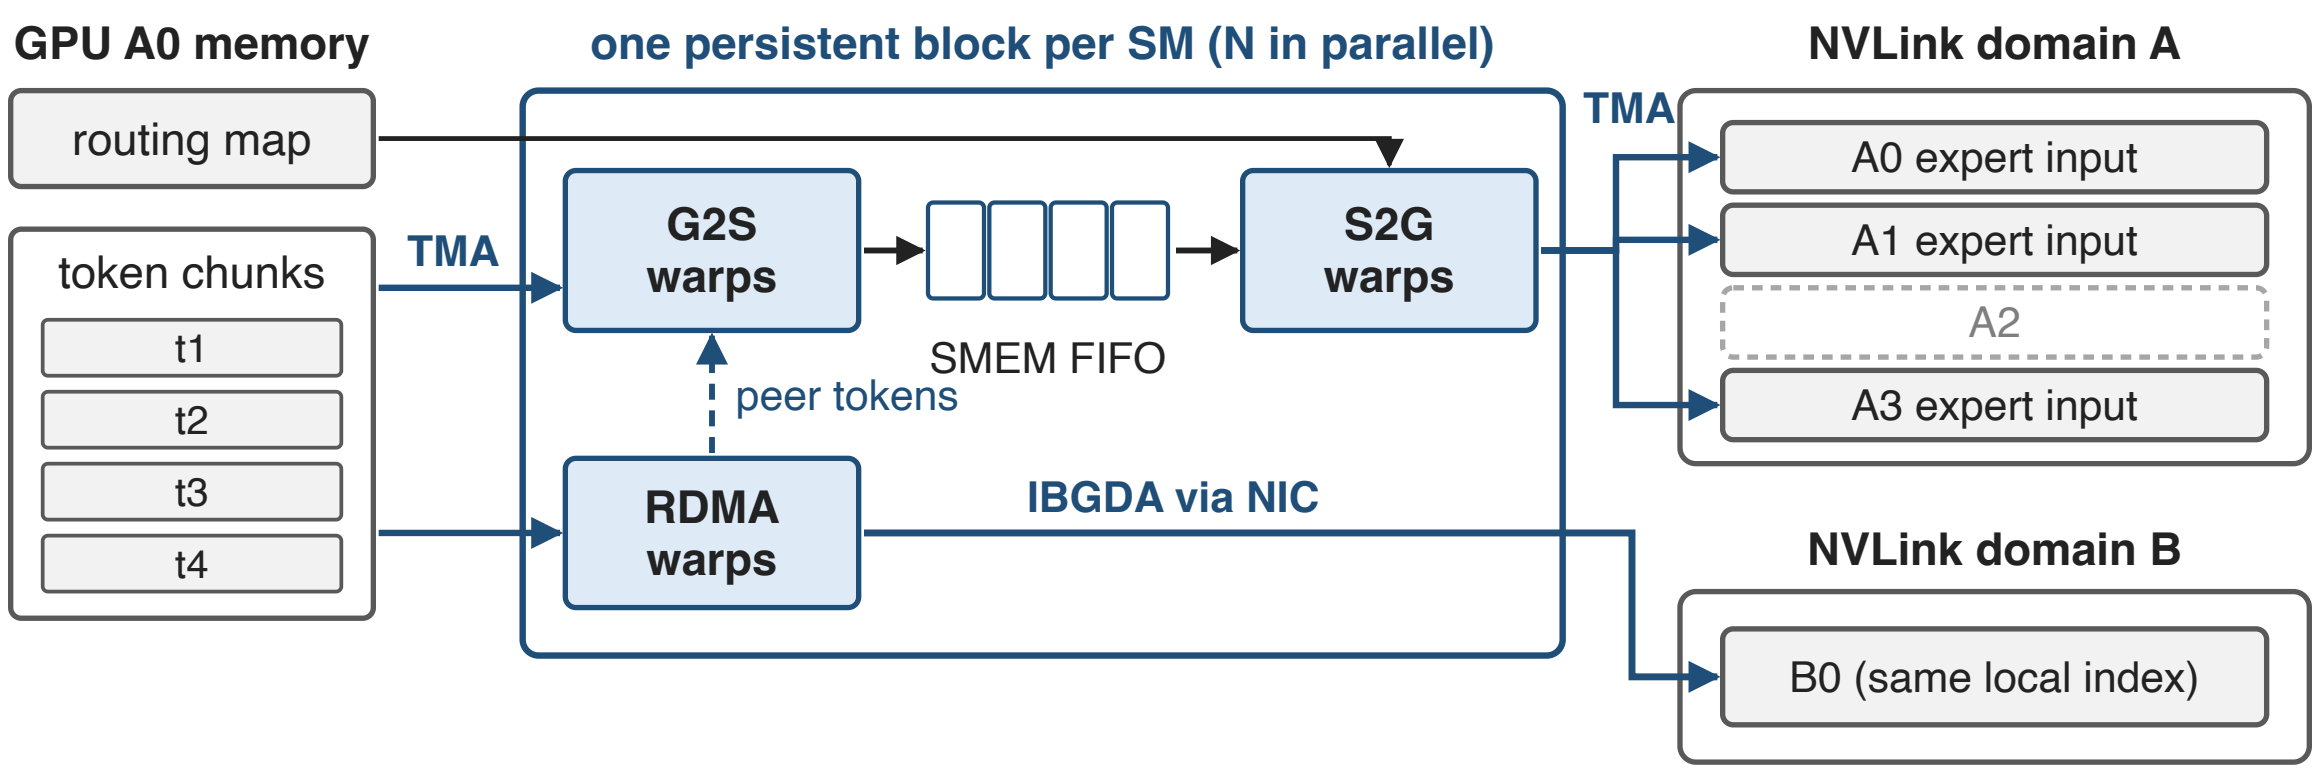
GPU A0 memory
routing map
token chunks
t1
t2
t3
t4
TMA
one persistent block per SM (N in parallel)
G2S warps
SMEM FIFO
S2G warps
peer tokens
RDMA warps
IBGDA via NIC
TMA
NVLink domain A
A0 expert input
A1 expert input
A2
A3 expert input
NVLink domain B
B0 (same local index)