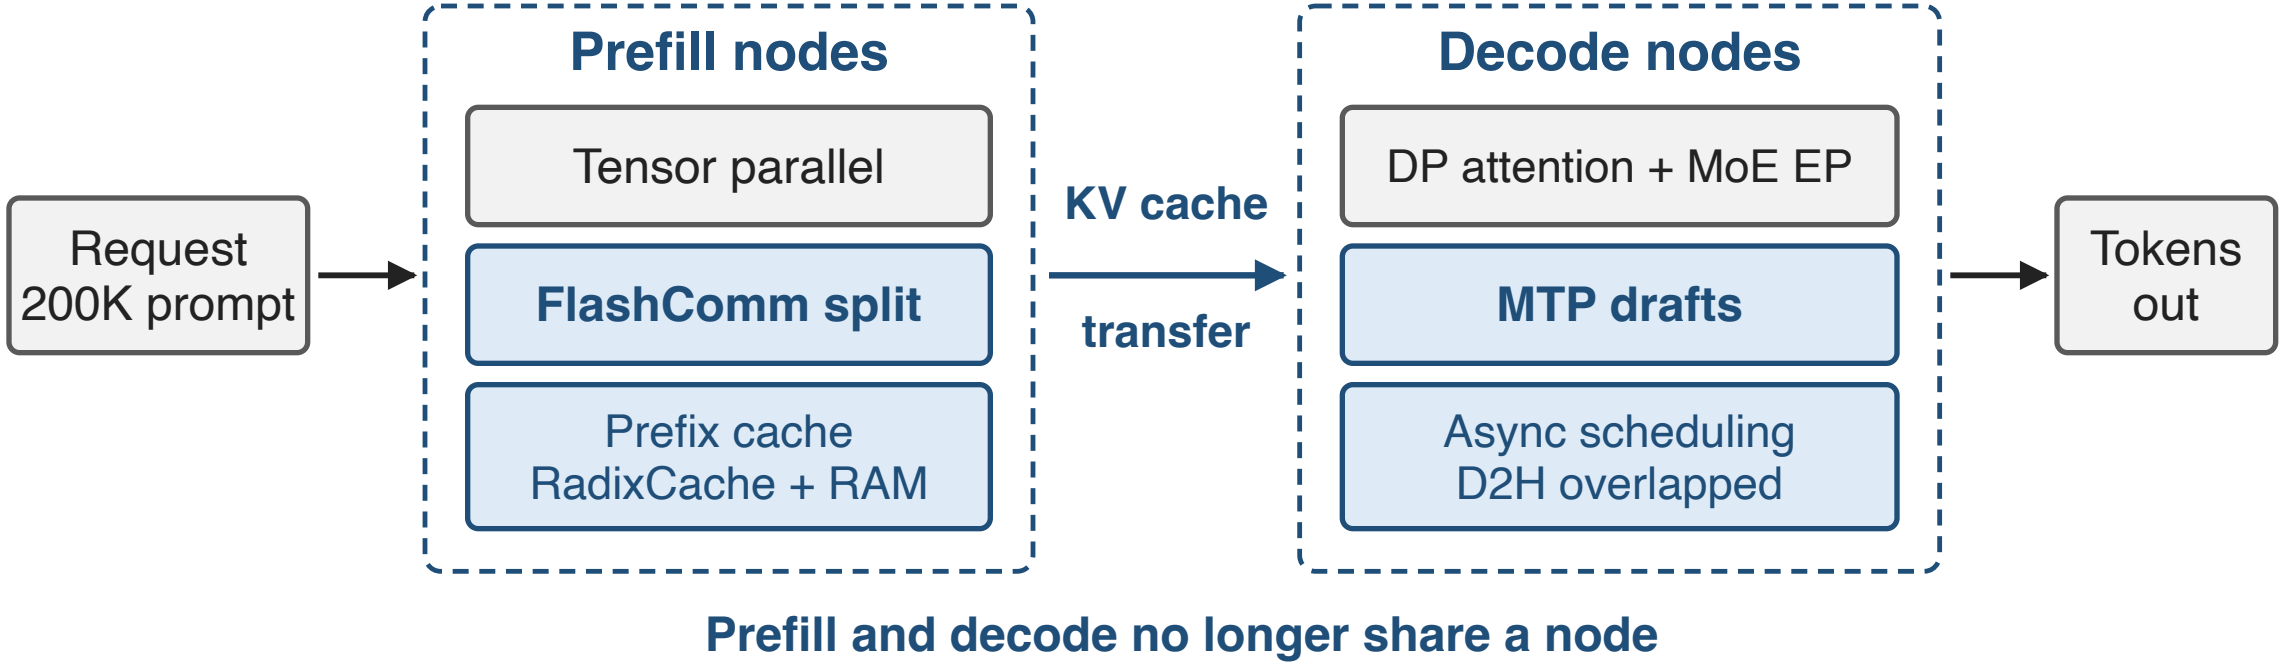
Request
200K prompt
Prefill nodes
Tensor parallel
FlashComm split
Prefix cache
RadixCache + RAM
KV cache
transfer
Decode nodes
DP attention + MoE EP
MTP drafts
Async scheduling
D2H overlapped
Tokens
out
Prefill and decode no longer share a node