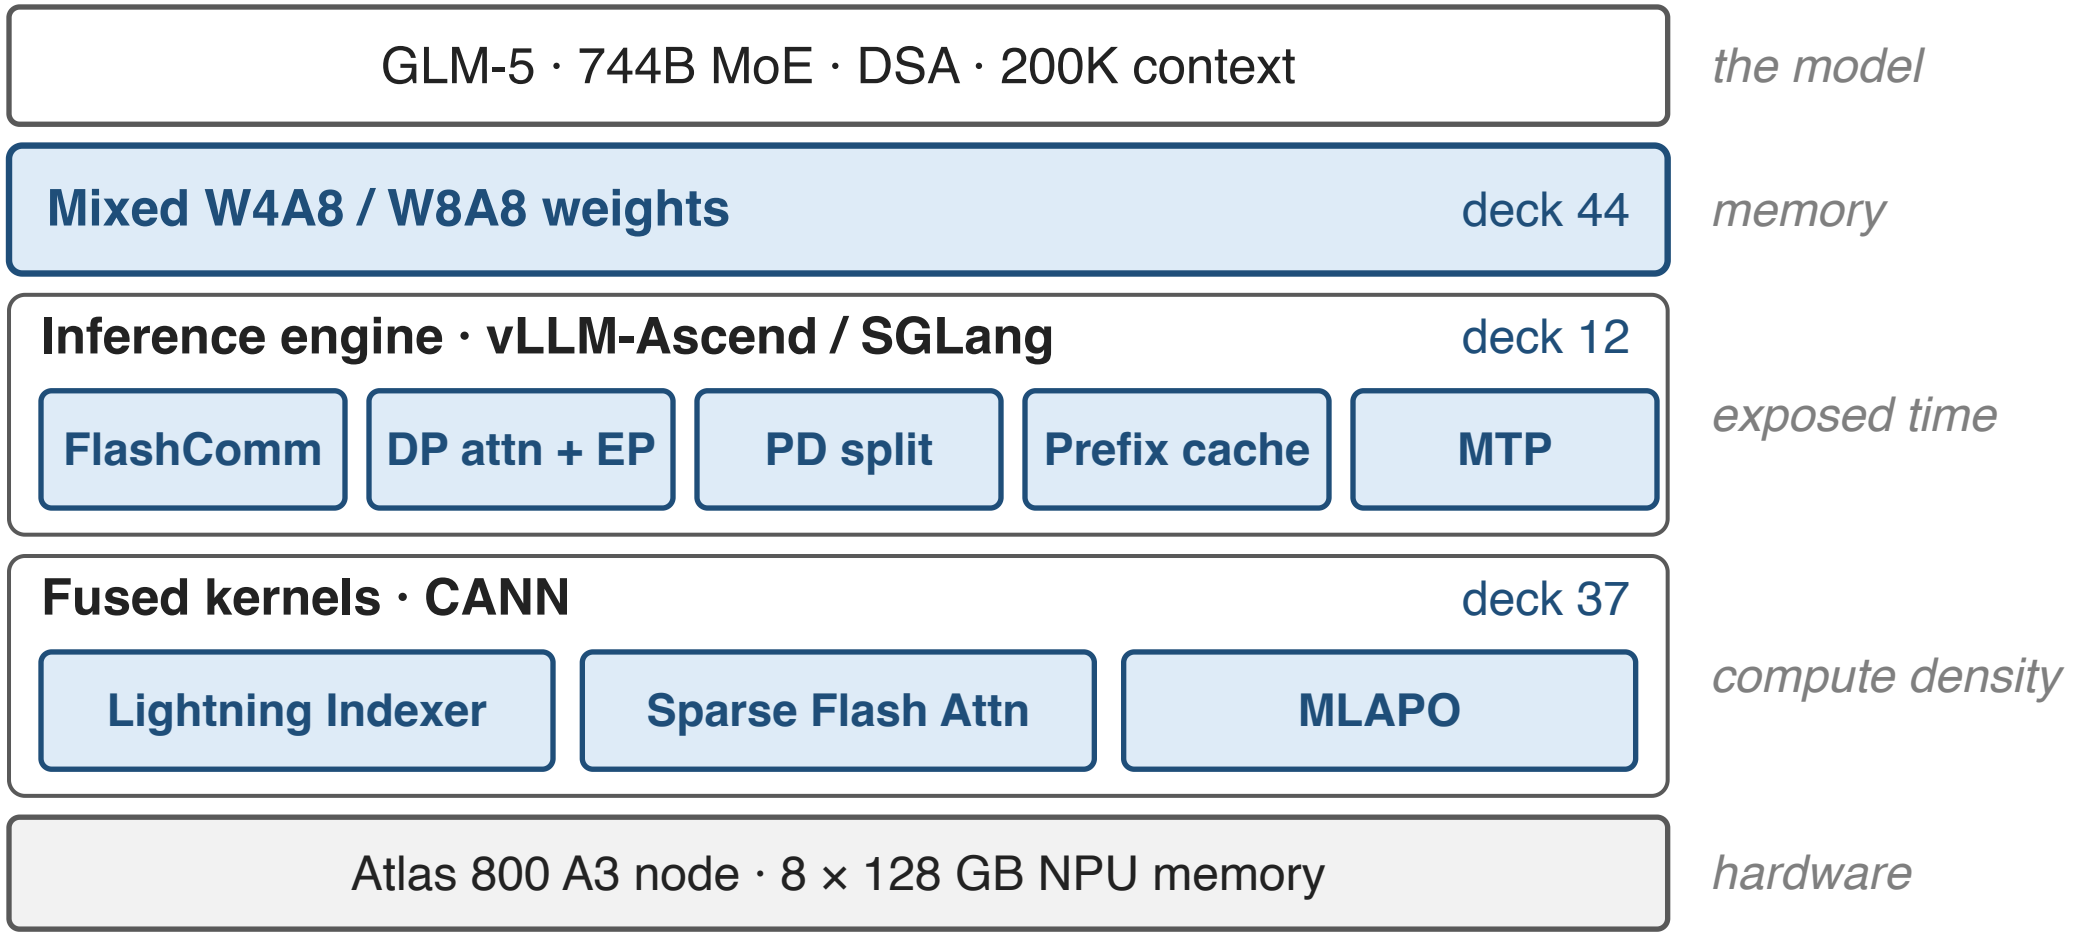
GLM-5 · 744B MoE · DSA · 200K context
the model
Mixed W4A8 / W8A8 weights
deck 44
memory
Inference engine · vLLM-Ascend / SGLang
deck 12
FlashComm
DP attn + EP
PD split
Prefix cache
MTP
exposed time
Fused kernels · CANN
deck 37
Lightning Indexer
Sparse Flash Attn
MLAPO
compute density
Atlas 800 A3 node · 8 × 128 GB NPU memory
hardware