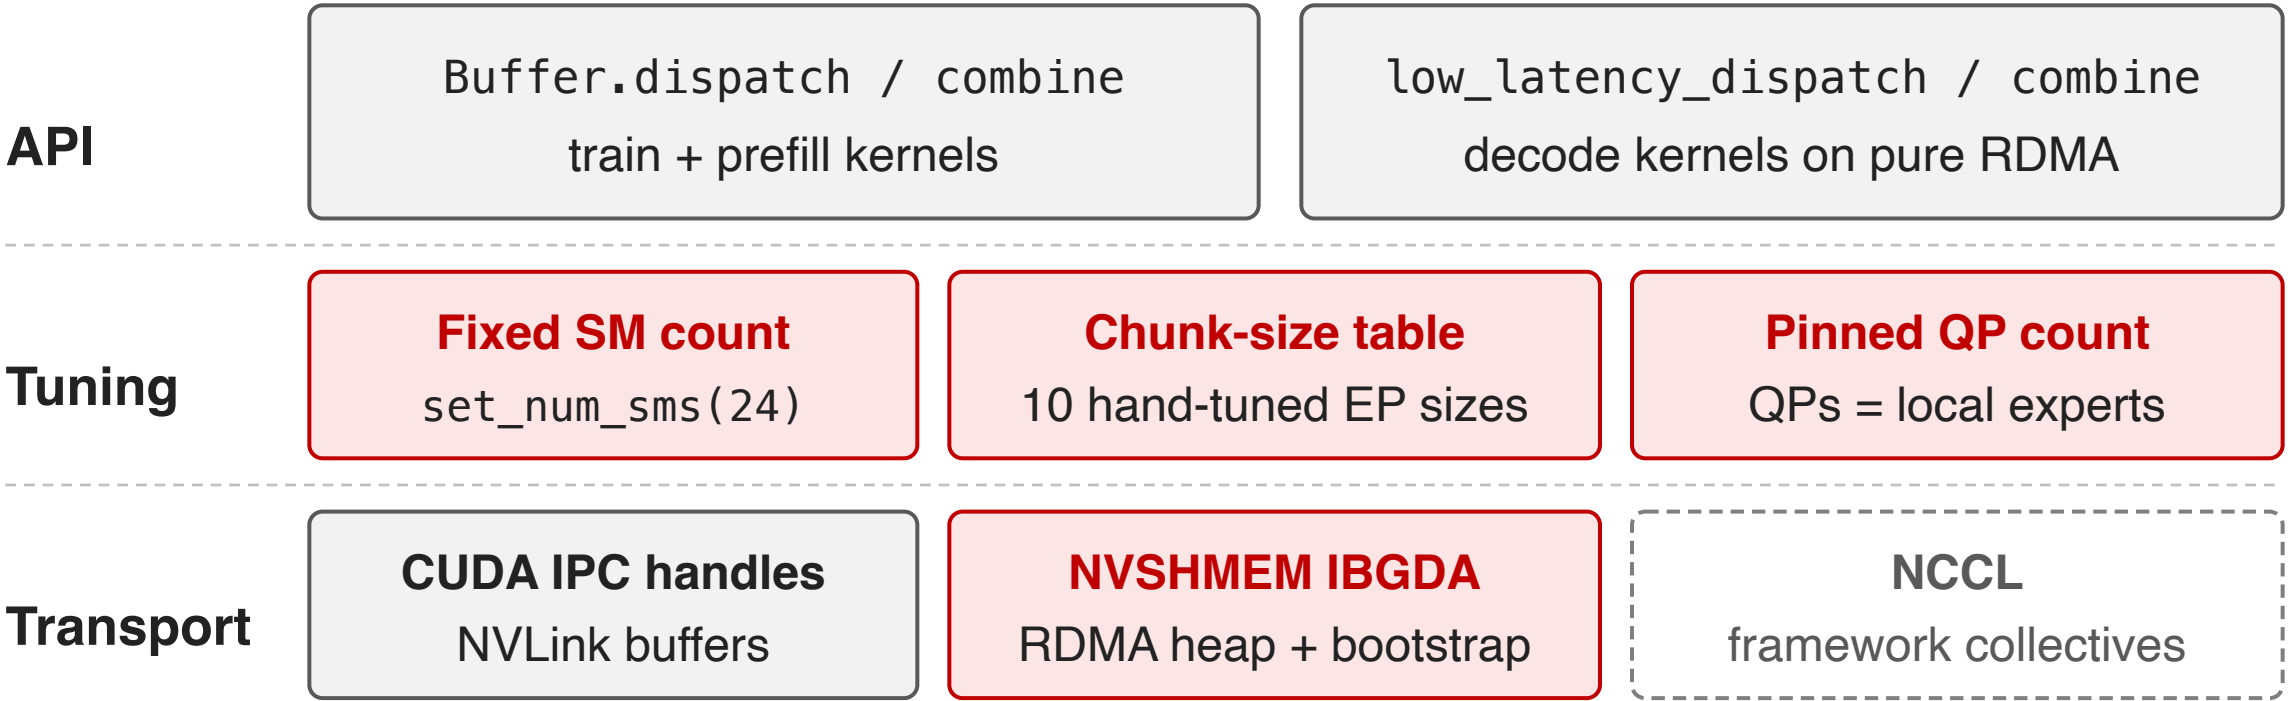
API
Buffer.dispatch / combine
train + prefill kernels
low_latency_dispatch / combine
decode kernels on pure RDMA
Tuning
Fixed SM count
set_num_sms(24)
Chunk-size table
10 hand-tuned EP sizes
Pinned QP count
QPs = local experts
Transport
CUDA IPC handles
NVLink buffers
NVSHMEM IBGDA
RDMA heap + bootstrap
NCCL
framework collectives