	A BLAS FP8 GEMM call	One DeepSeek-V4 MoE step
Scaling	one scale per tensor	one scale per 1x128 tile along K
Shapes	M, N, K fixed on the host	tokens per expert known on the GPU
Operands	one dtype for A and B	FP8 activations x FP4 weights
Around it	matmul only	EP dispatch, SwiGLU, combine