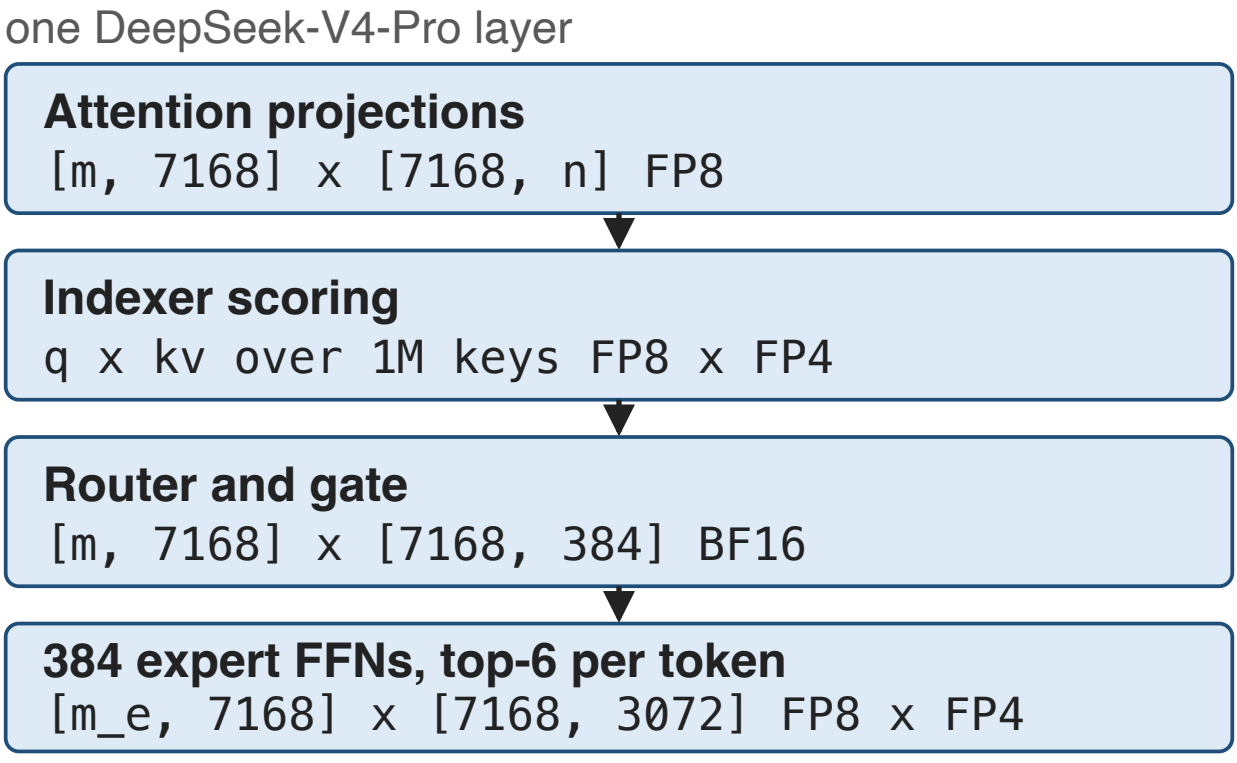
one DeepSeek-V4-Pro layer
Attention projections
[m, 7168] x [7168, n] FP8
Indexer scoring
q x kv over 1M keys FP8 x FP4
Router and gate
[m, 7168] x [7168, 384] BF16
384 expert FFNs, top-6 per token
[m_e, 7168] x [7168, 3072] FP8 x FP4