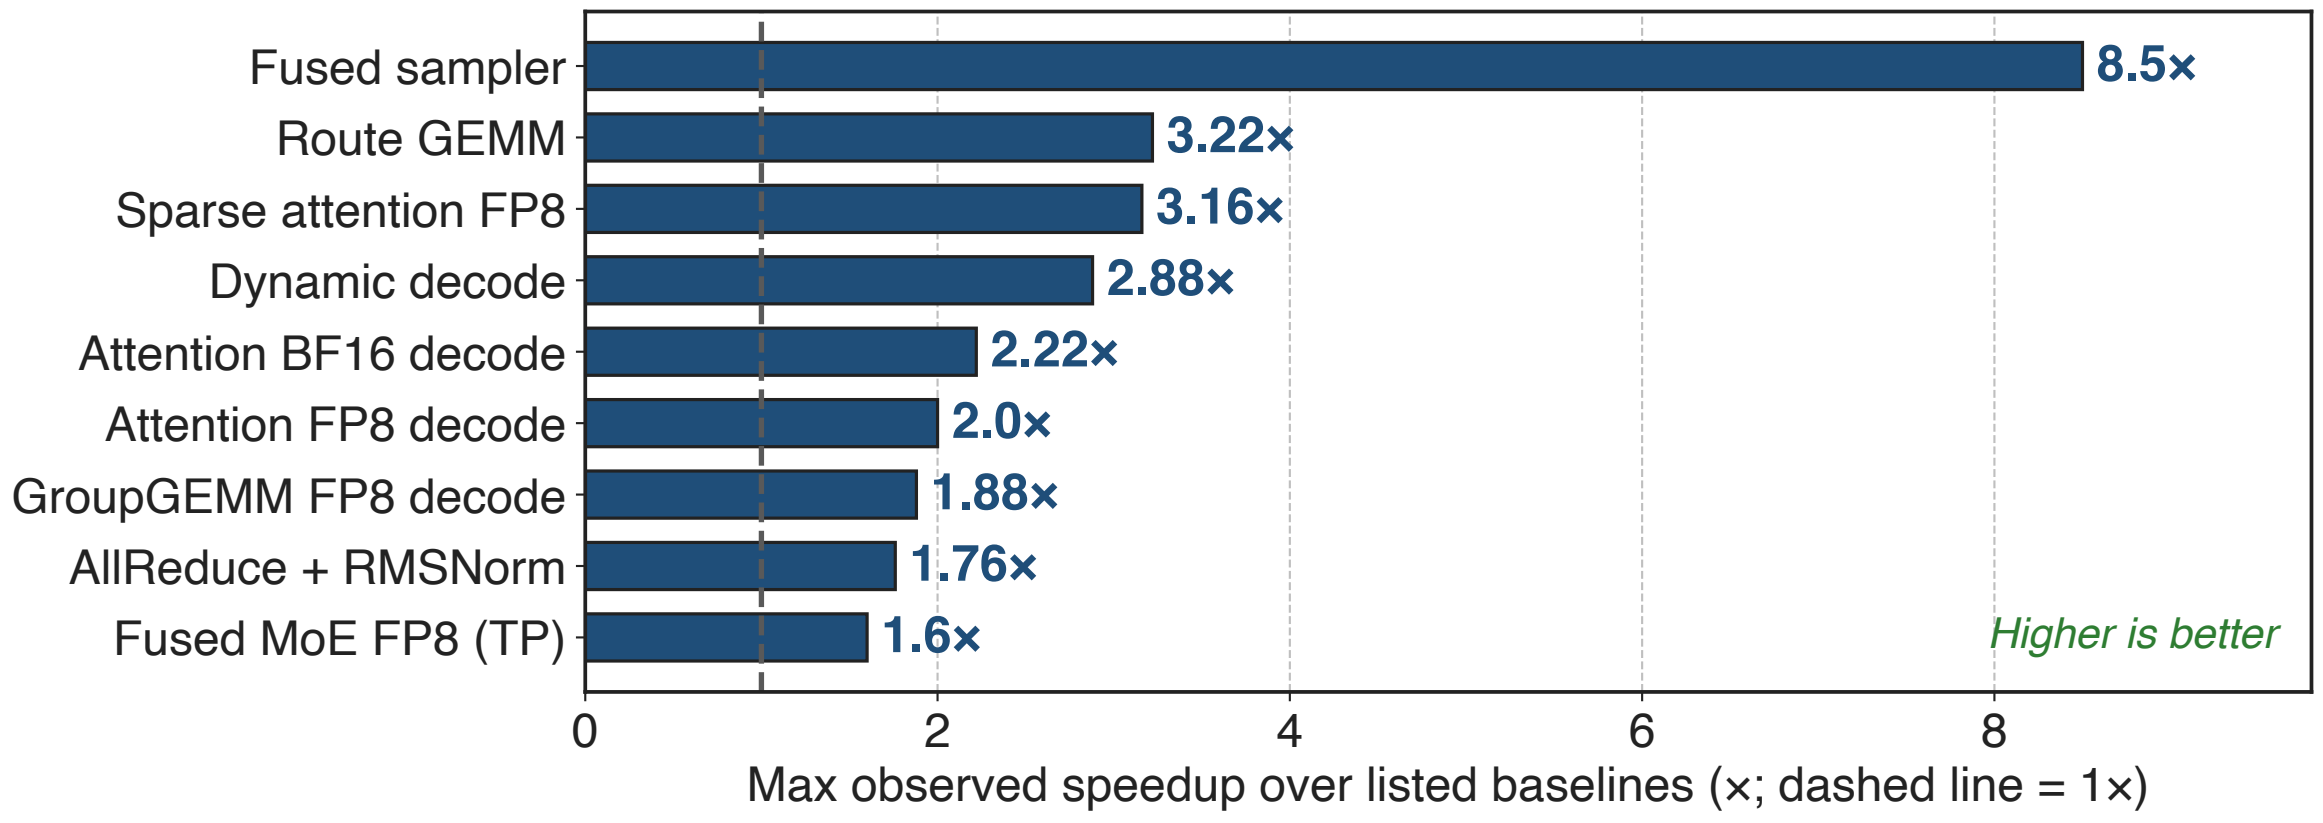
Fused sampler
Route GEMM
Sparse attention FP8
Dynamic decode
Attention BF16 decode
Attention FP8 decode
GroupGEMM FP8 decode
AllReduce + RMSNorm
Fused MoE FP8 (TP)
8.5×
3.22×
3.16×
2.88×
2.22×
2.0×
1.88×
1.76×
1.6×
Higher is better
0
2
4
6
8
Max observed speedup over listed baselines (×; dashed line = 1×)