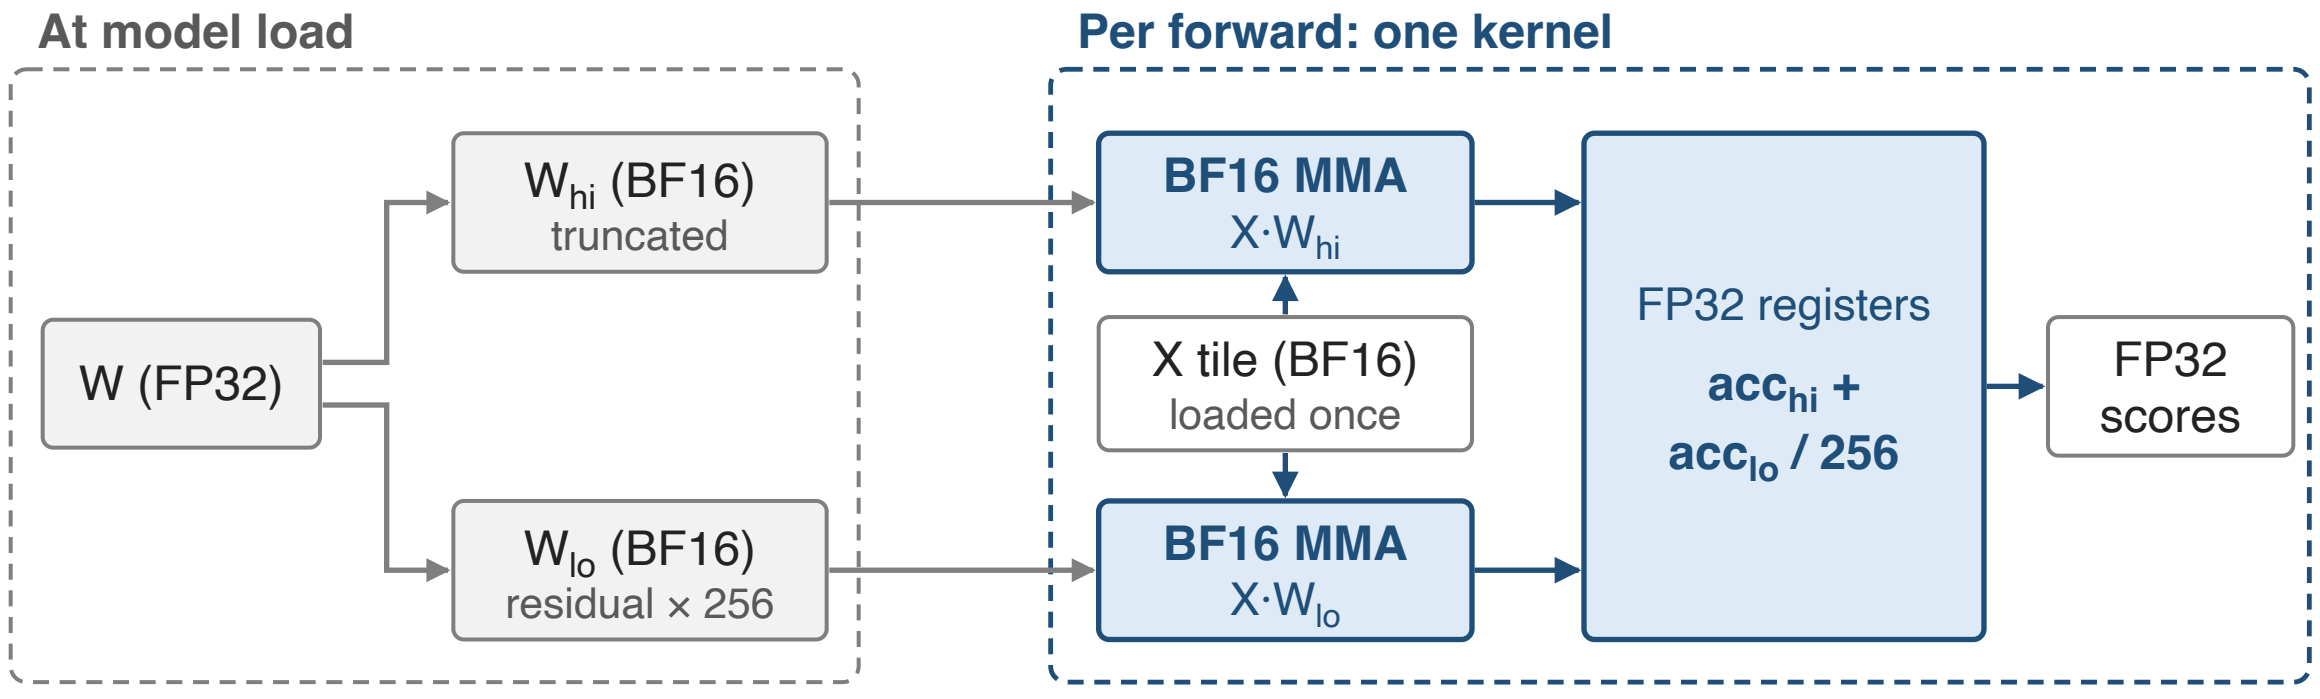
At model load
W (FP32)
W_{hi} (BF16)
truncated
W_{lo} (BF16)
residual × 256
Per forward: one kernel
BF16 MMA
$X \cdot W_{hi}$
X tile (BF16)
loaded once
BF16 MMA
$X \cdot W_{lo}$
FP32 registers
acc_{hi} +
acc_{lo} / 256
FP32
scores