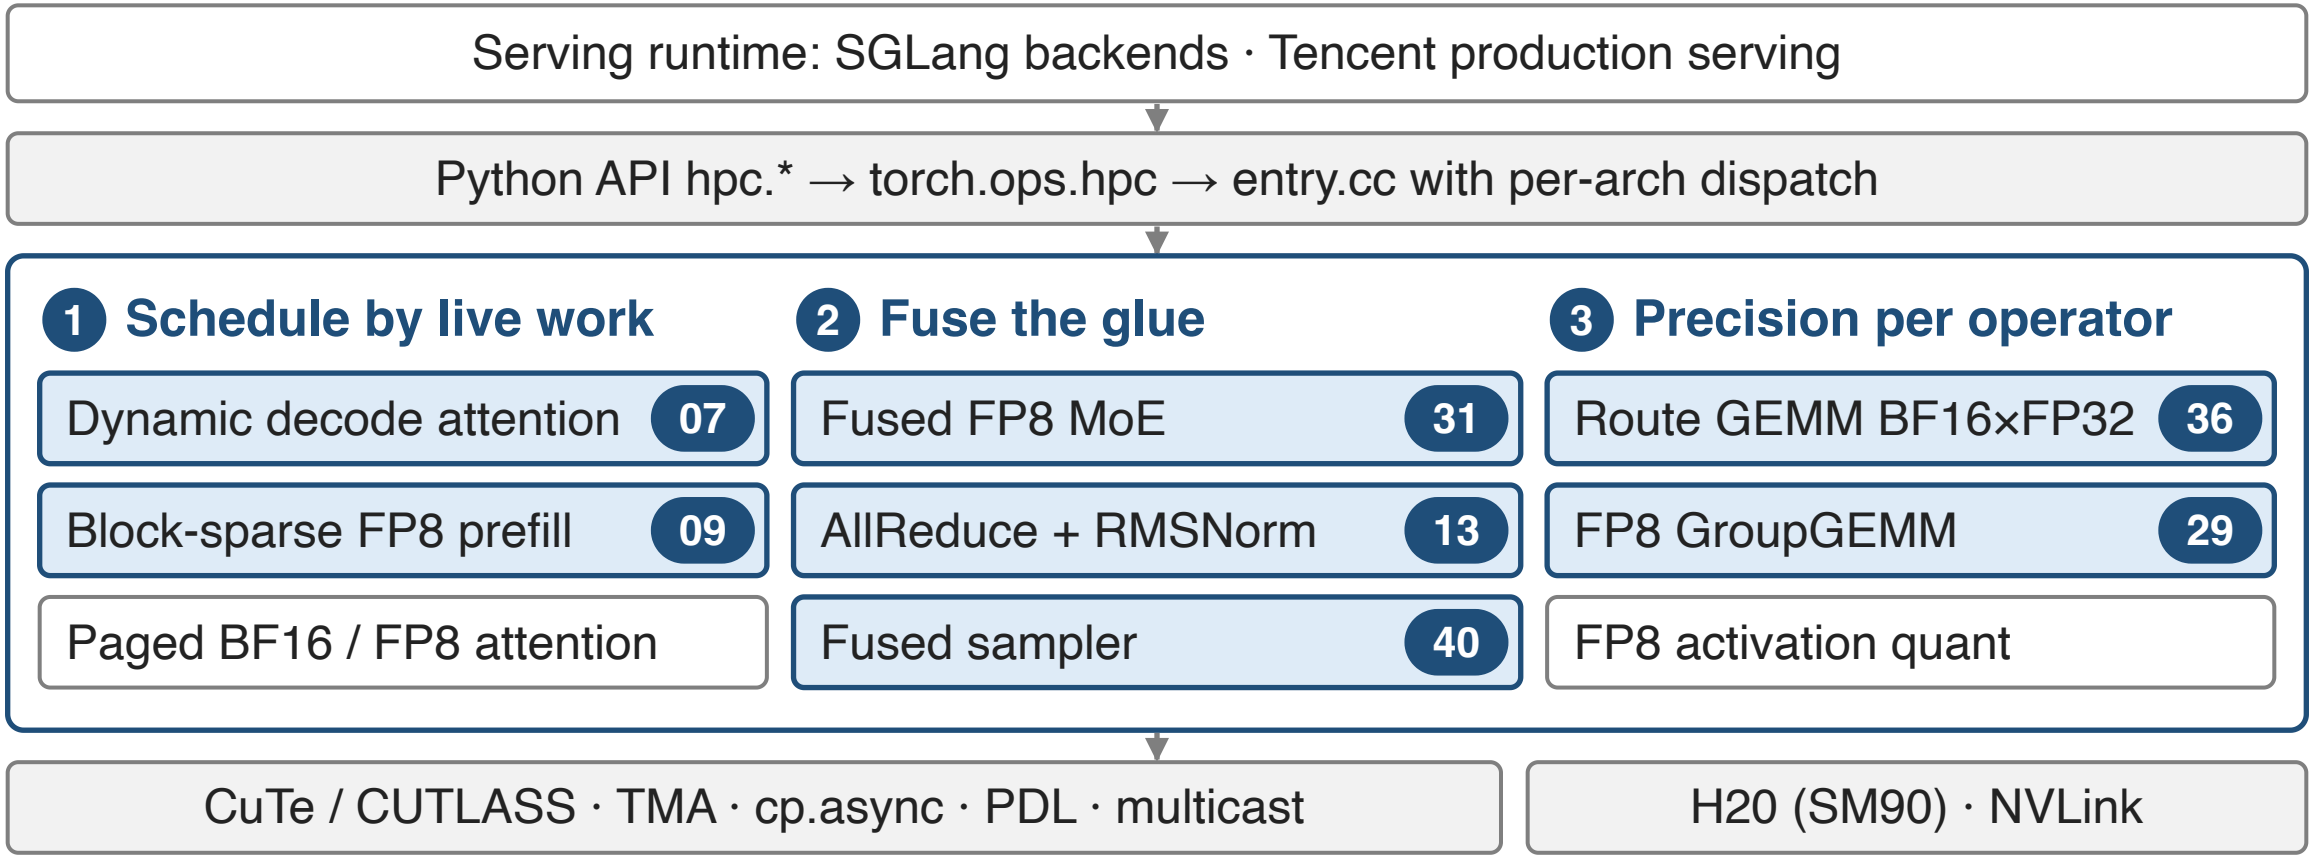
Serving runtime: SGLang backends · Tencent production serving
Python API hpc.* → torch.ops.hpc → entry.cc with per-arch dispatch
1 Schedule by live work
Dynamic decode attention 07
Block-sparse FP8 prefill 09
Paged BF16 / FP8 attention
2 Fuse the glue
Fused FP8 MoE 31
AllReduce + RMSNorm 13
Fused sampler 40
3 Precision per operator
Route GEMM BF16×FP32 36
FP8 GroupGEMM 29
FP8 activation quant
CuTe / CUTLASS · TMA · cp.async · PDL · multicast
H20 (SM90) · NVLink