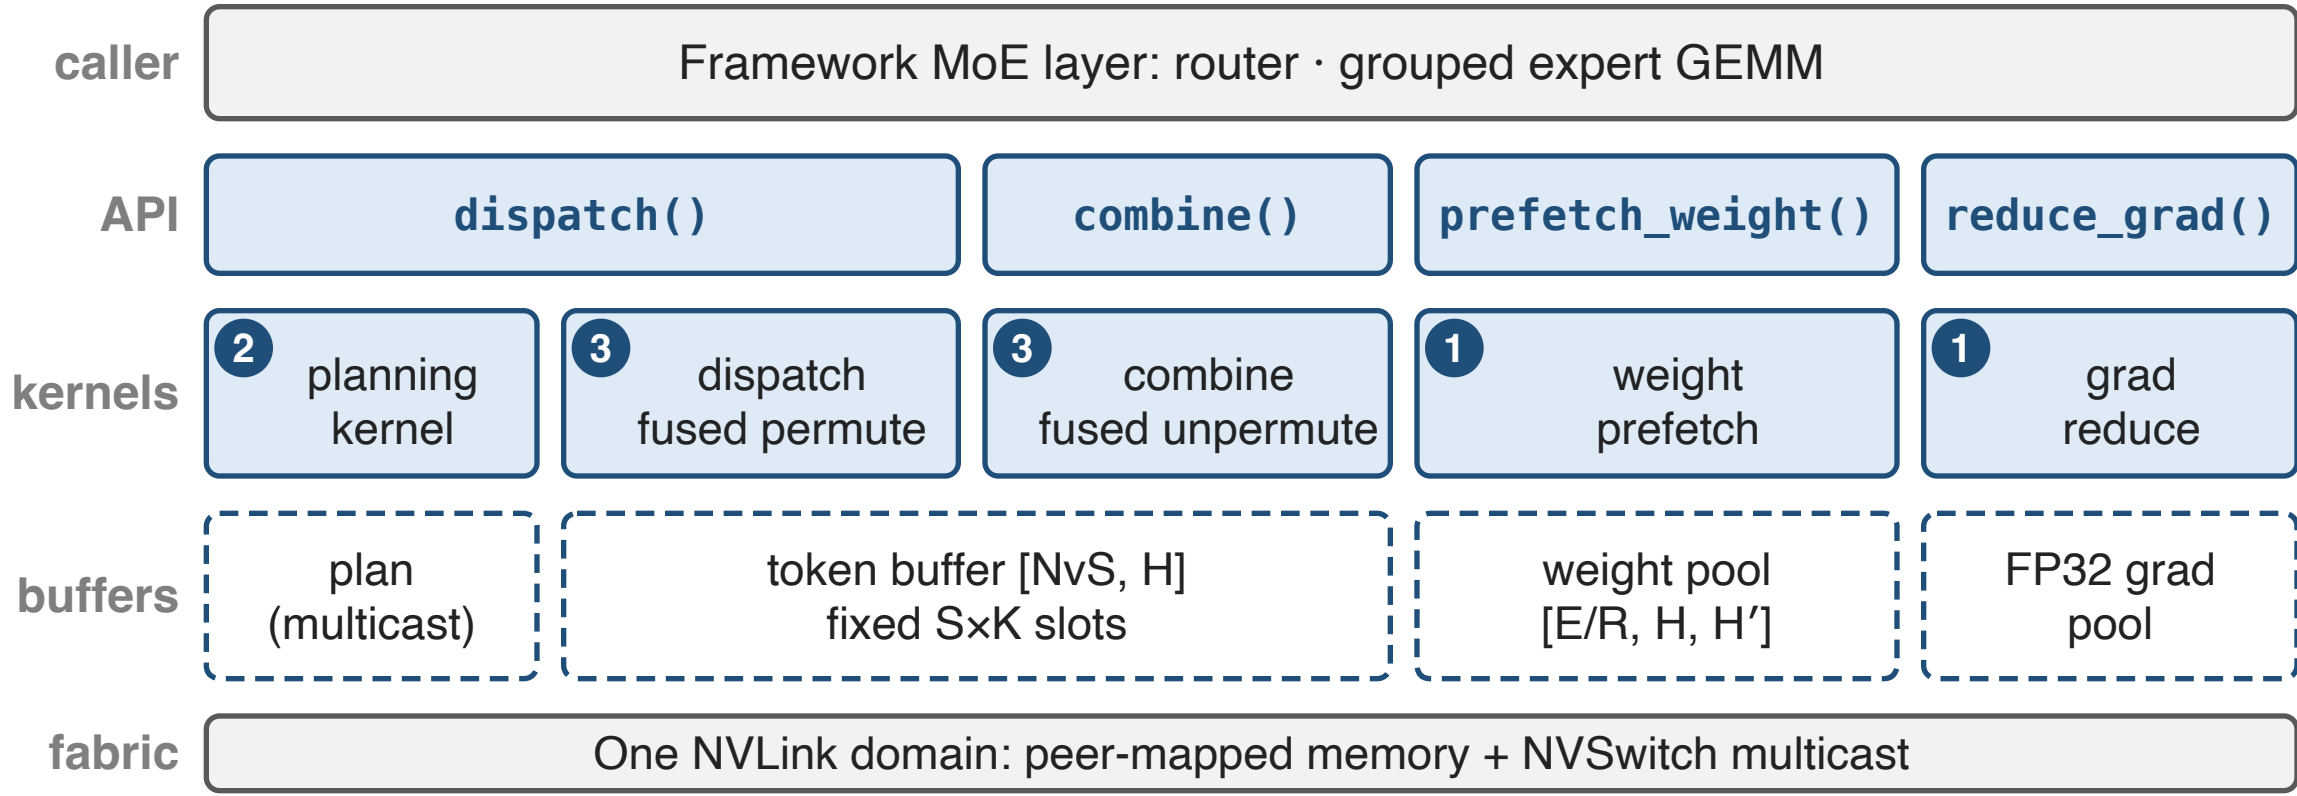
caller
Framework MoE layer: router · grouped expert GEMM
API
dispatch()
combine()
prefetch_weight()
reduce_grad()
kernels
2
planning kernel
3
dispatch fused permute
3
combine fused unpermute
1
weight prefetch
1
grad reduce
buffers
plan (multicast)
token buffer [NvS, H] fixed S×K slots
weight pool [E/R, H, H′]
FP32 grad pool
fabric
One NVLink domain: peer-mapped memory + NVSwitch multicast