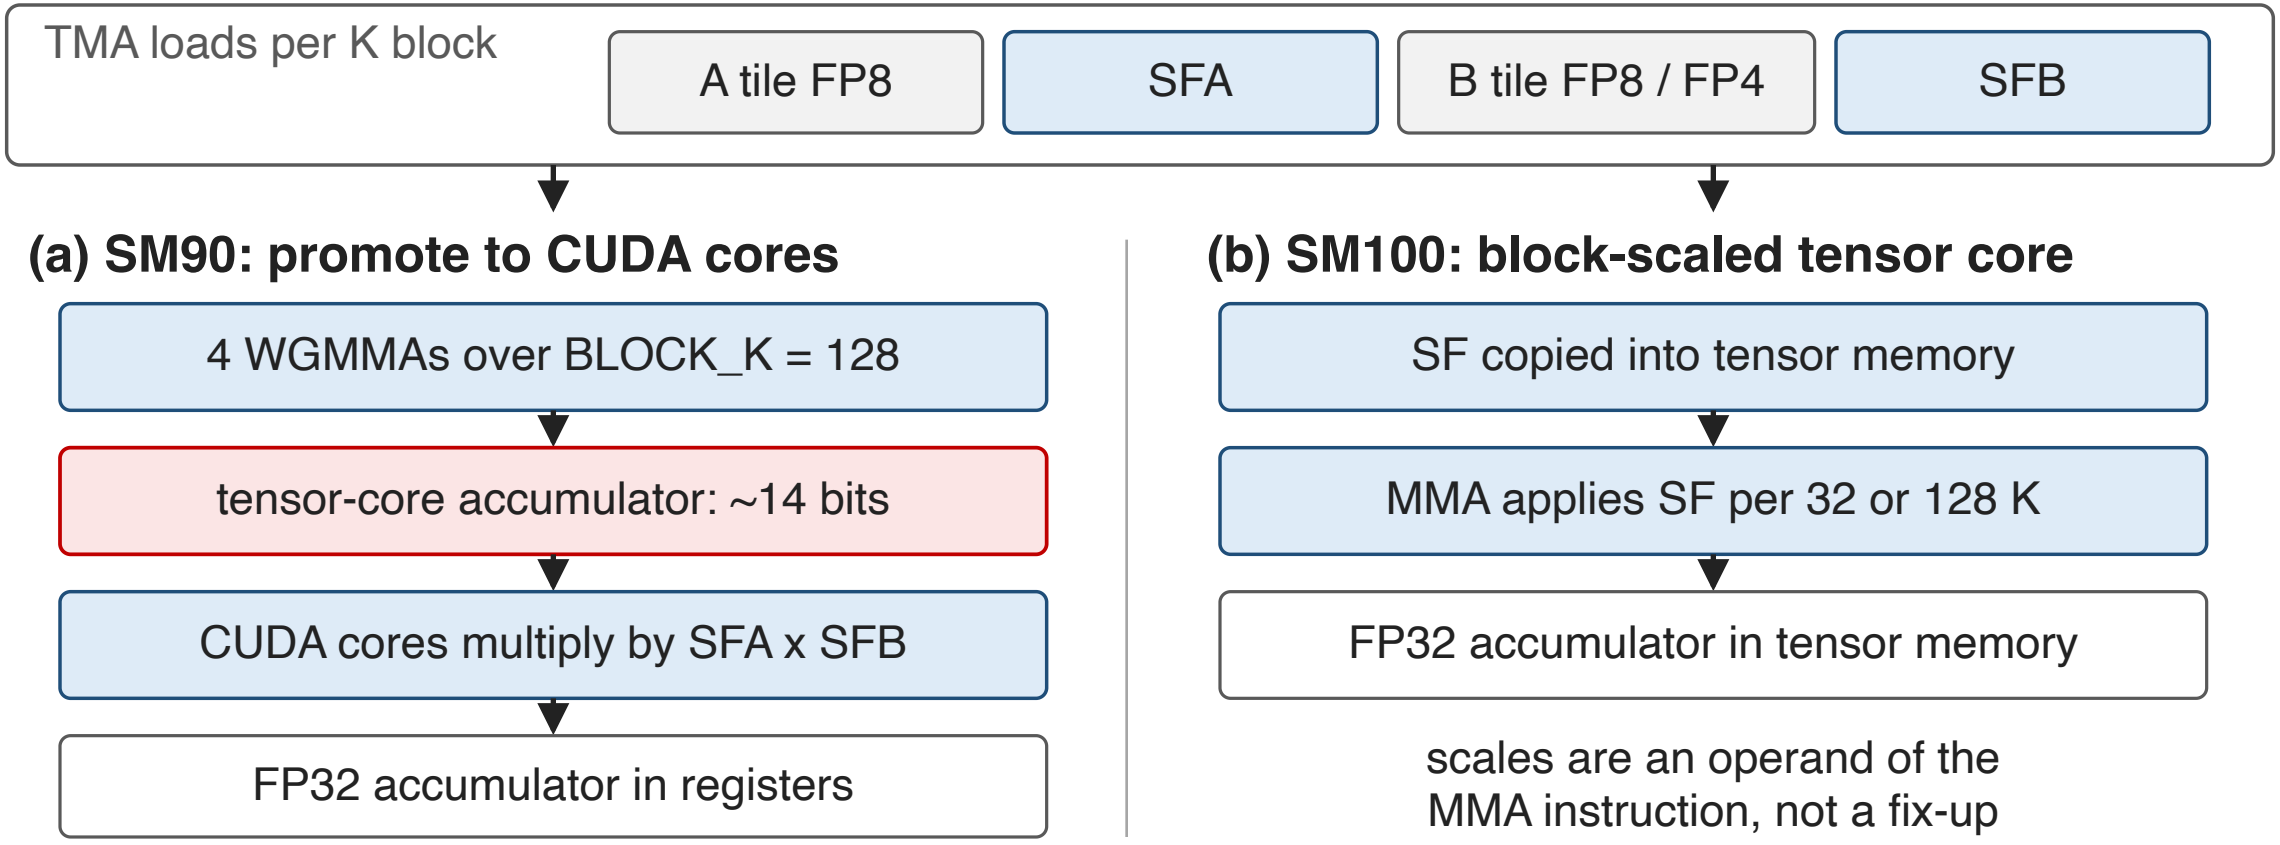
TMA loads per K block
A tile FP8
SFA
B tile FP8 / FP4
SFB
(a) SM90: promote to CUDA cores
4 WGMMAs over BLOCK_K = 128
tensor-core accumulator: ~14 bits
CUDA cores multiply by SFA x SFB
FP32 accumulator in registers
(b) SM100: block-scaled tensor core
SF copied into tensor memory
MMA applies SF per 32 or 128 K
FP32 accumulator in tensor memory
scales are an operand of the
MMA instruction, not a fix-up