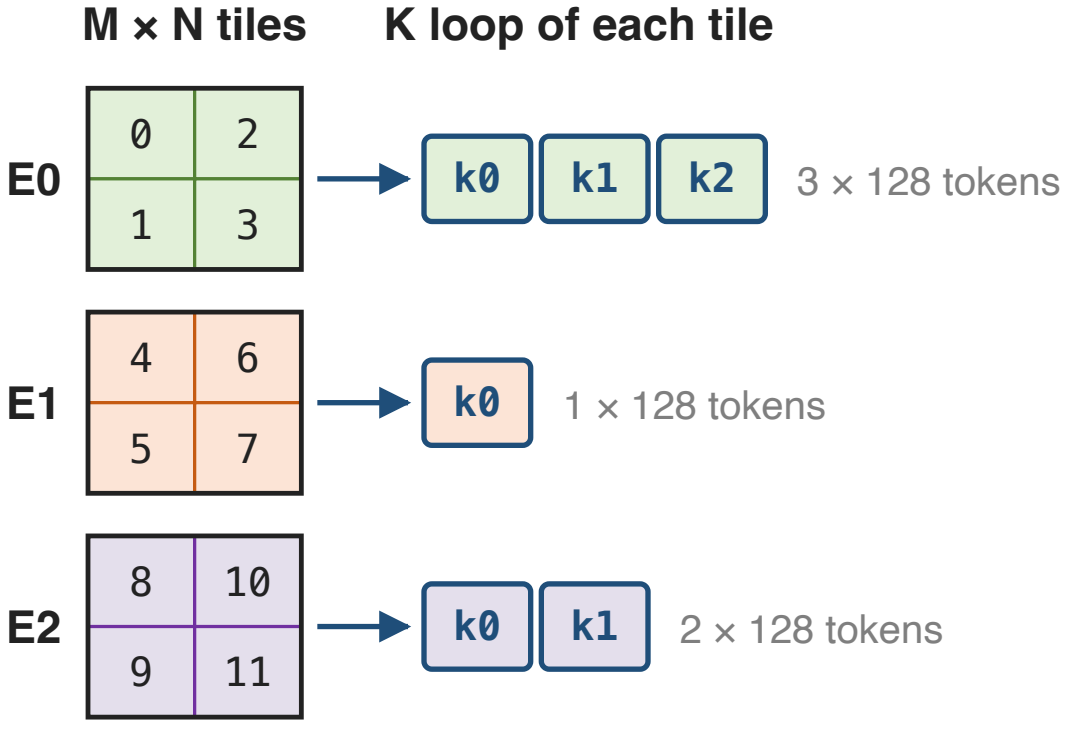

one CTA per SM (persistent)

```
for t in range(sm, G*tiles, SMs):        all experts
  g, m, n = split(t)                     same grid
  acc = C[g][m, n]                       fused accum
  for kb in k_blocks(g):                 own tokens
    acc += A[kb,m].T @ B[kb,n]           FP8 MMA
  D[g][m, n] = acc                       FP32 store
```

k_blocks(g) from grouped_layout · start aligned to 128