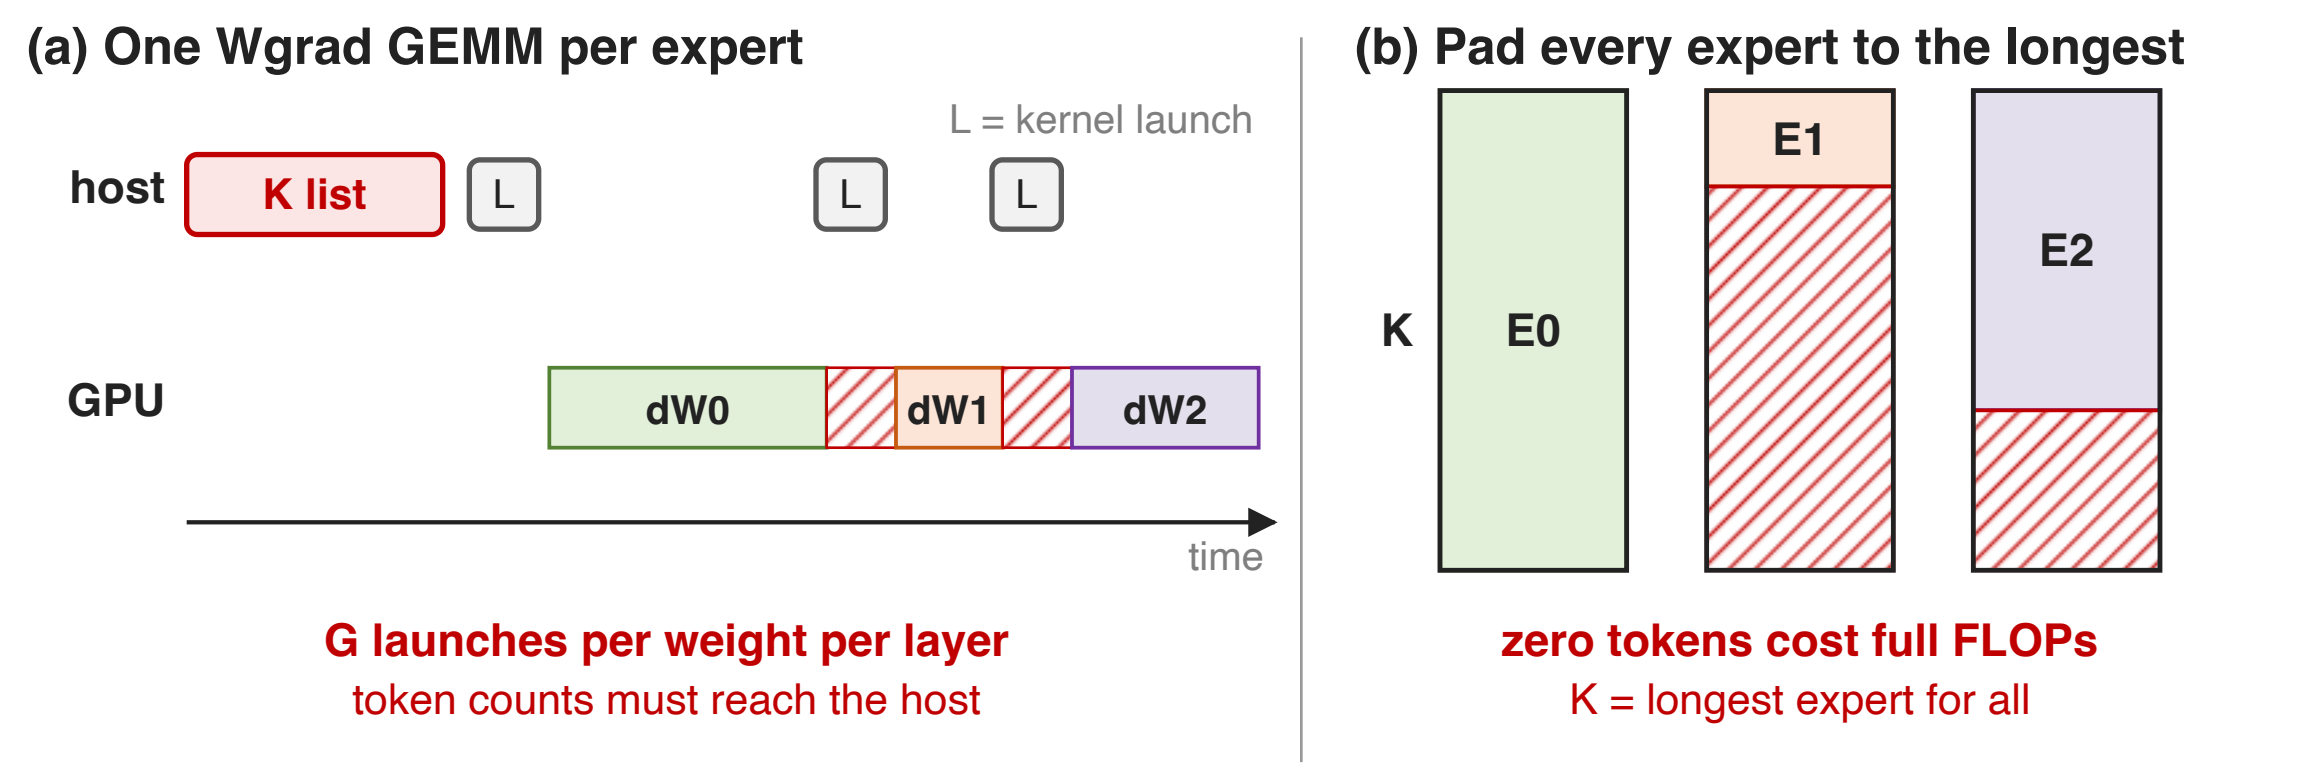

(a) One Wgrad GEMM per expert
L = kernel launch
host
K list
L
L
L
GPU
dW0
dW1
dW2
time
G launches per weight per layer
token counts must reach the host
(b) Pad every expert to the longest
K
E0
E1
E2
zero tokens cost full FLOPs
K = longest expert for all