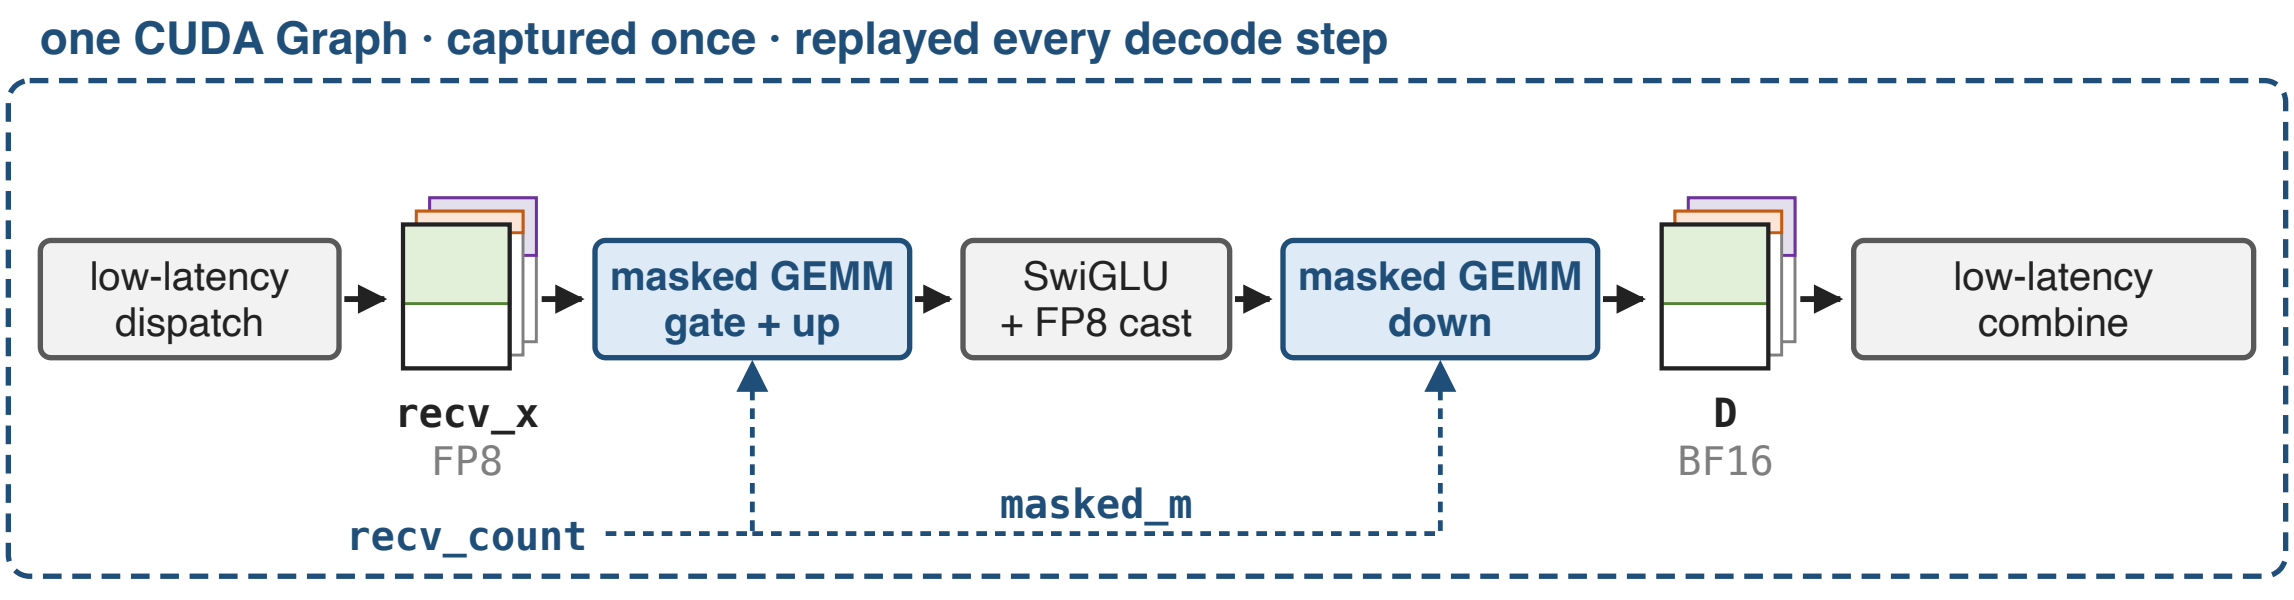
one CUDA Graph · captured once · replayed every decode step
low-latency dispatch
recv_x
FP8
masked GEMM gate + up
SwiGLU + FP8 cast
masked GEMM down
D
BF16
low-latency combine
recv_count
masked_m
slot shape [local experts, T × R, 7168]
T = max tokens per rank · R = EP ranks