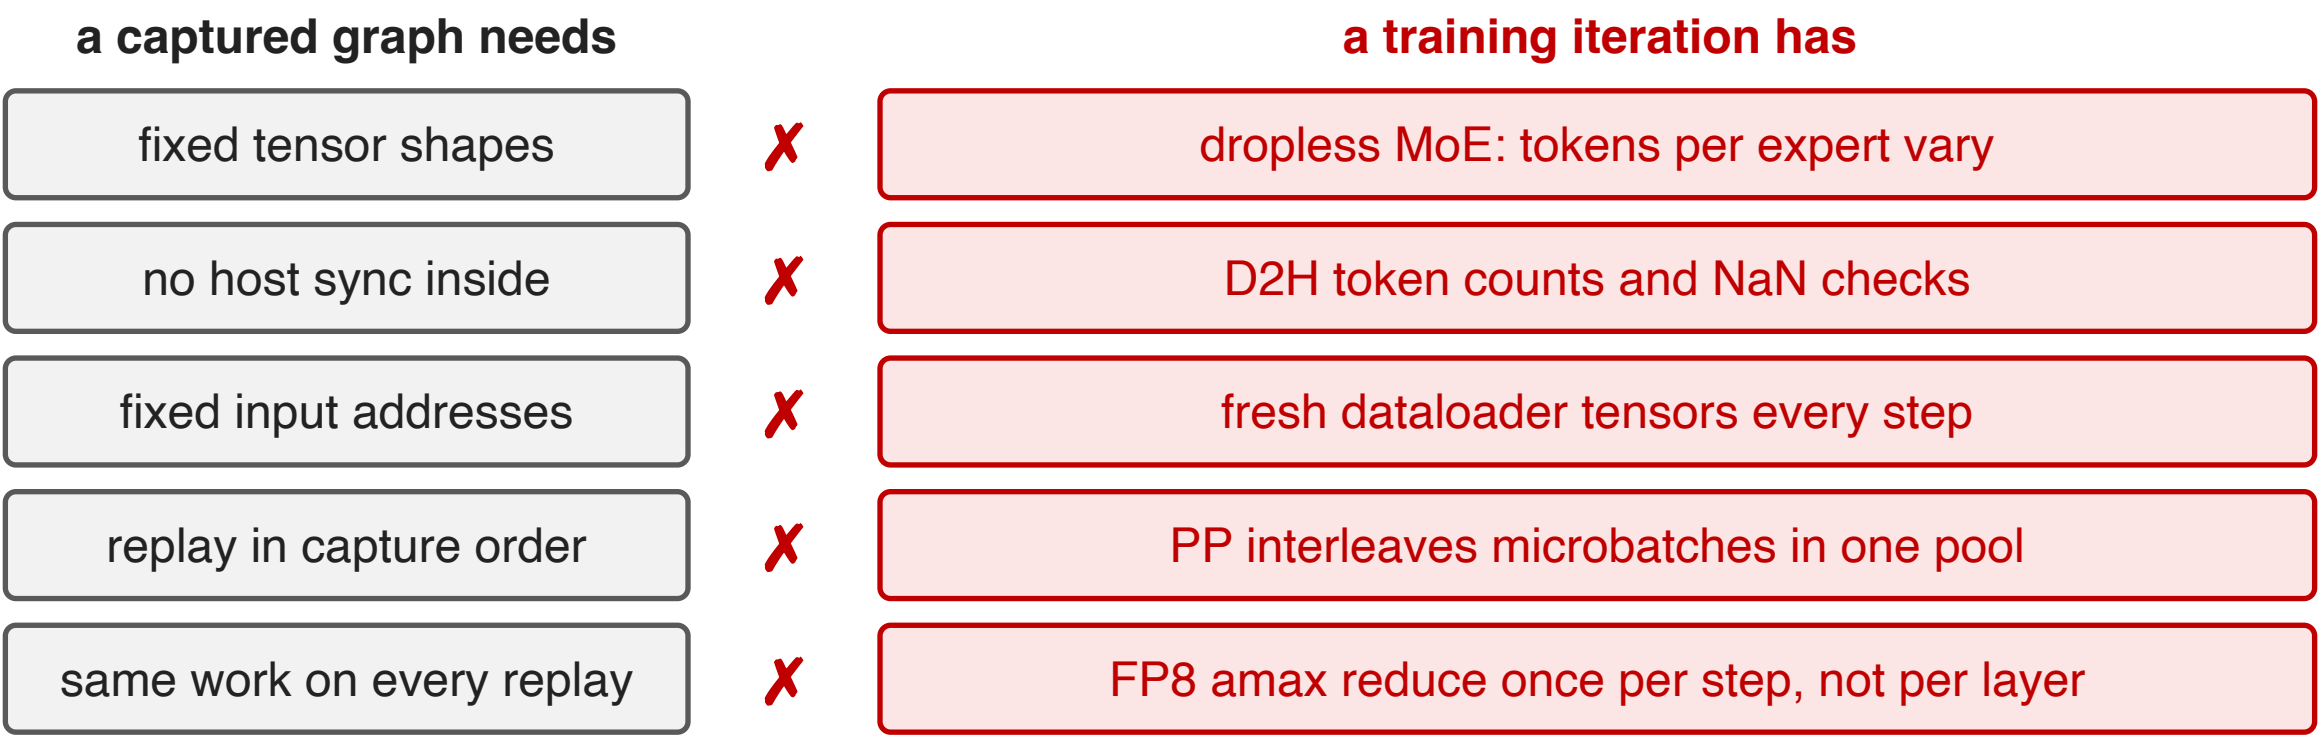
a captured graph needs
a training iteration has
fixed tensor shapes
dropless MoE: tokens per expert vary
no host sync inside
D2H token counts and NaN checks
fixed input addresses
fresh dataloader tensors every step
replay in capture order
PP interleaves microbatches in one pool
same work on every replay
FP8 amax reduce once per step, not per layer