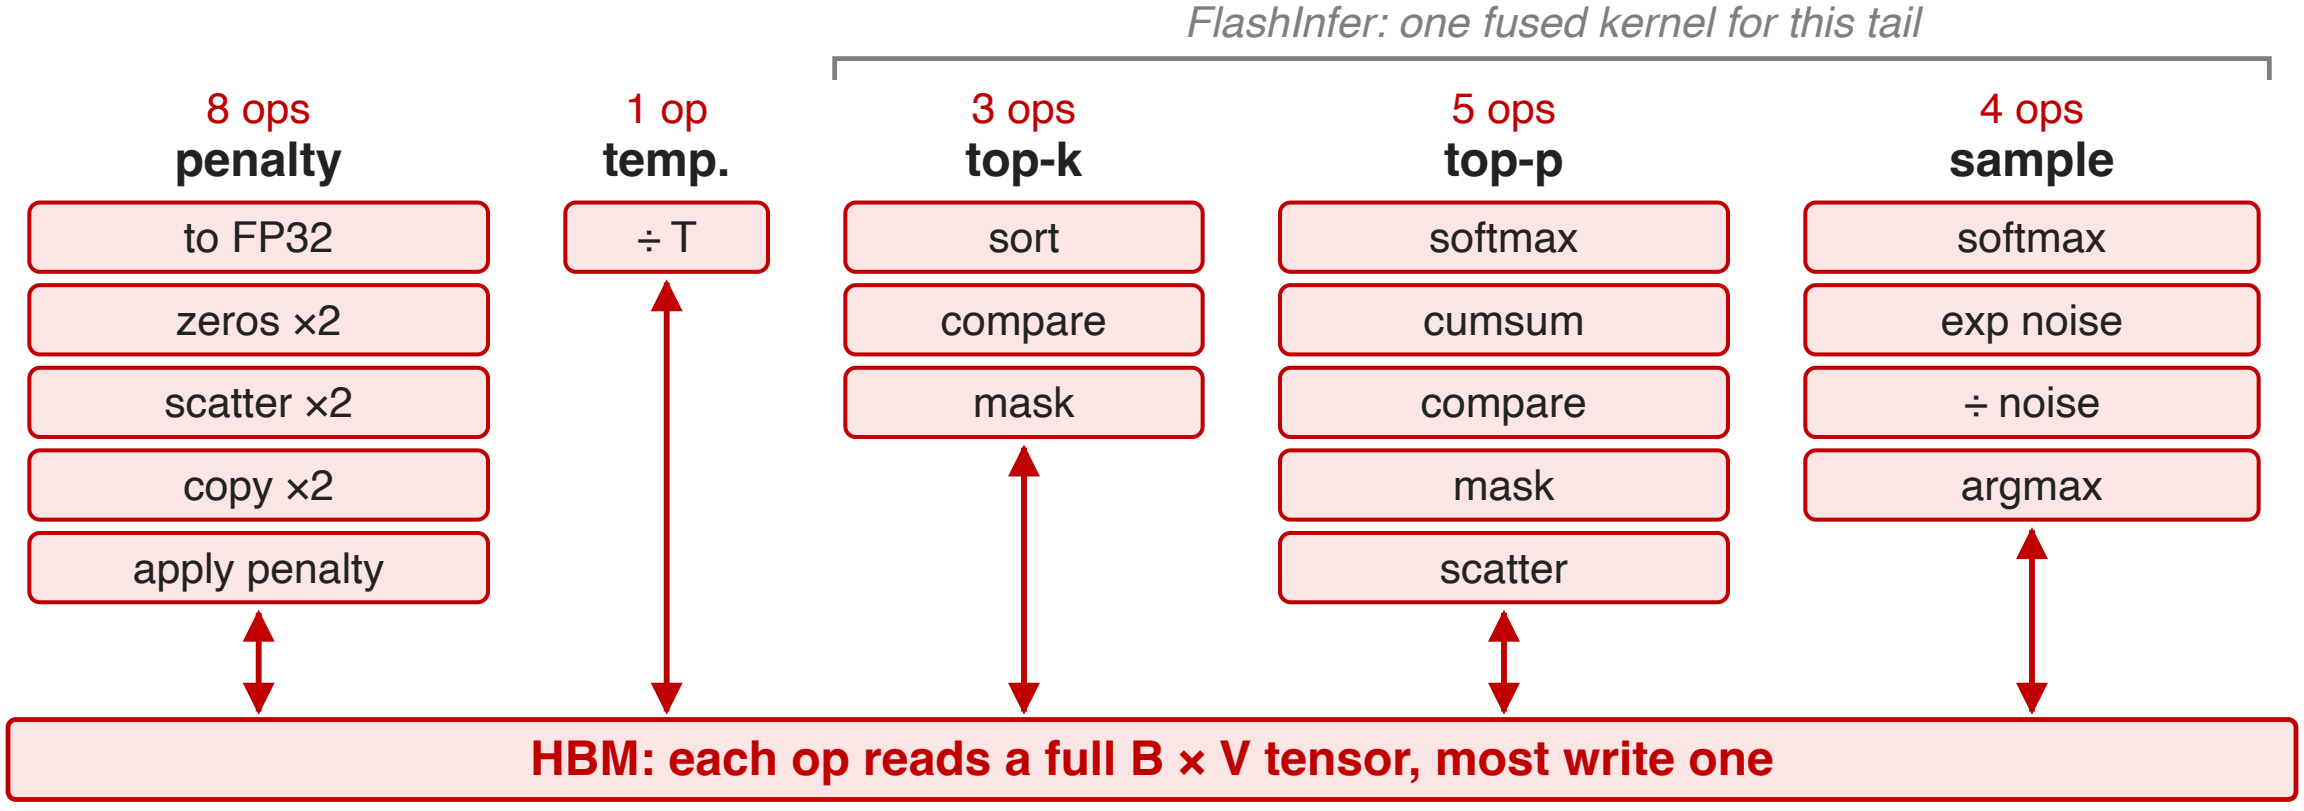
FlashInfer: one fused kernel for this tail
8 ops
penalty
to FP32
zeros ×2
scatter ×2
copy ×2
apply penalty
1 op
temp.
÷ T
3 ops
top-k
sort
compare
mask
5 ops
top-p
softmax
cumsum
compare
mask
scatter
4 ops
sample
softmax
exp noise
÷ noise
argmax
HBM: each op reads a full B × V tensor, most write one