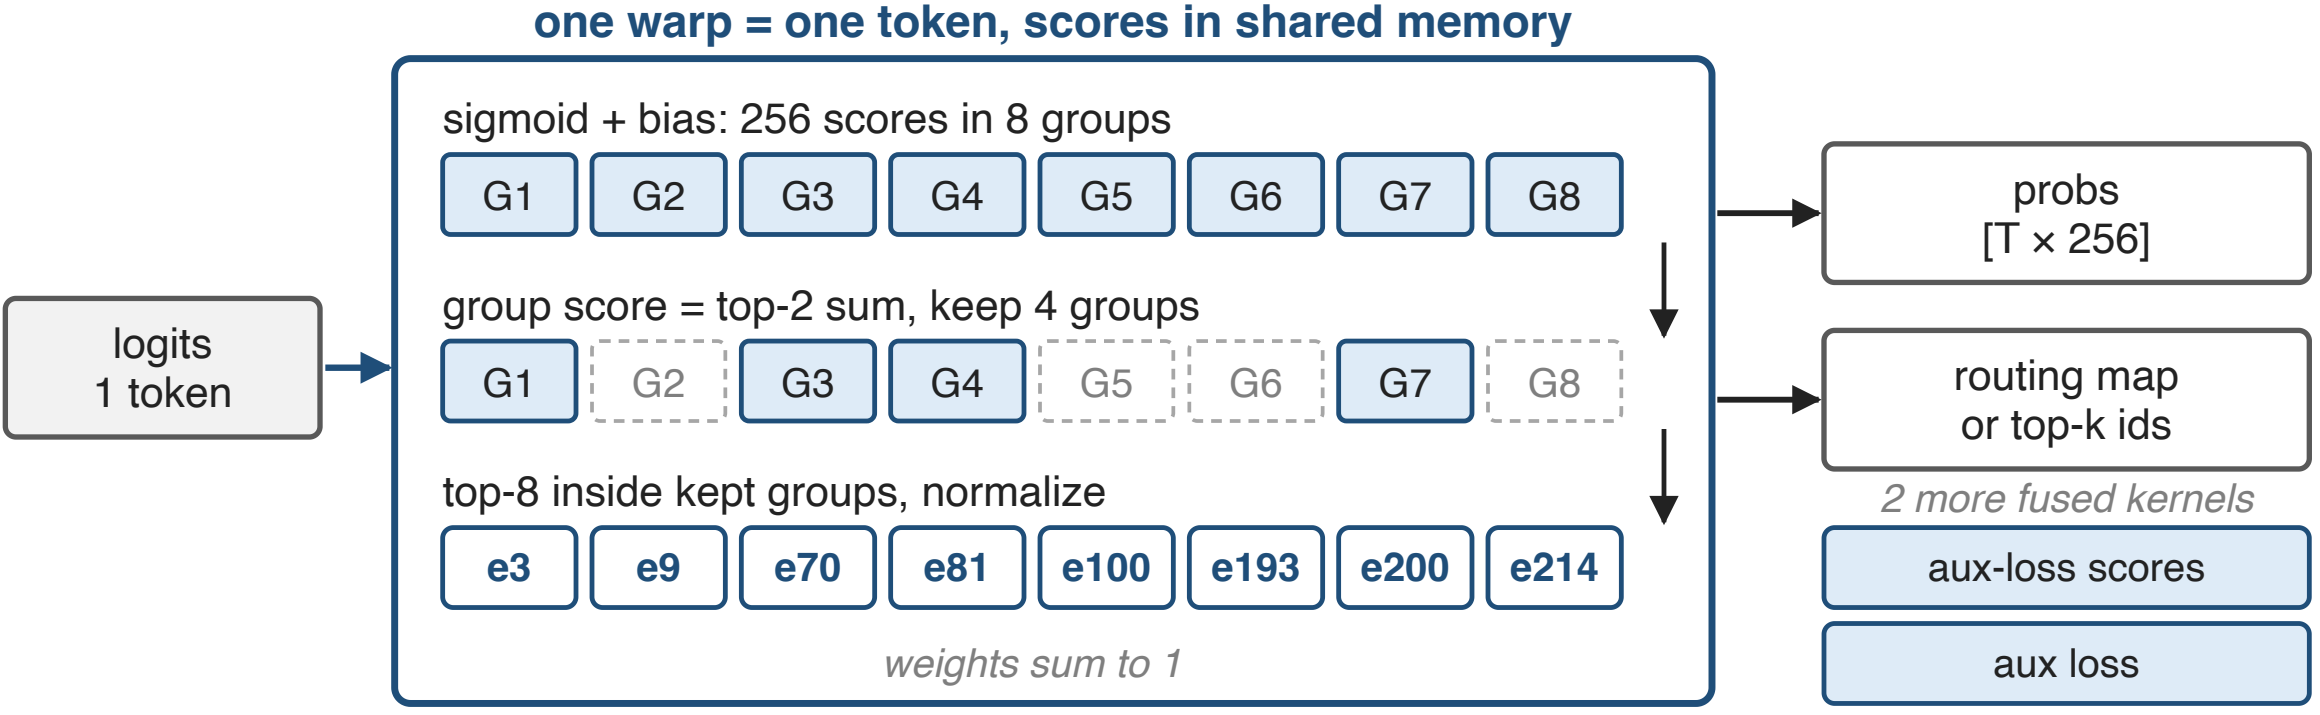
one warp = one token, scores in shared memory
logits
1 token
sigmoid + bias: 256 scores in 8 groups
G1
G2
G3
G4
G5
G6
G7
G8
group score = top-2 sum, keep 4 groups
G1
G2
G3
G4
G5
G6
G7
G8
top-8 inside kept groups, normalize
e3
e9
e70
e81
e100
e193
e200
e214
weights sum to 1
probs
[T × 256]
routing map
or top-k ids
2 more fused kernels
aux-loss scores
aux loss