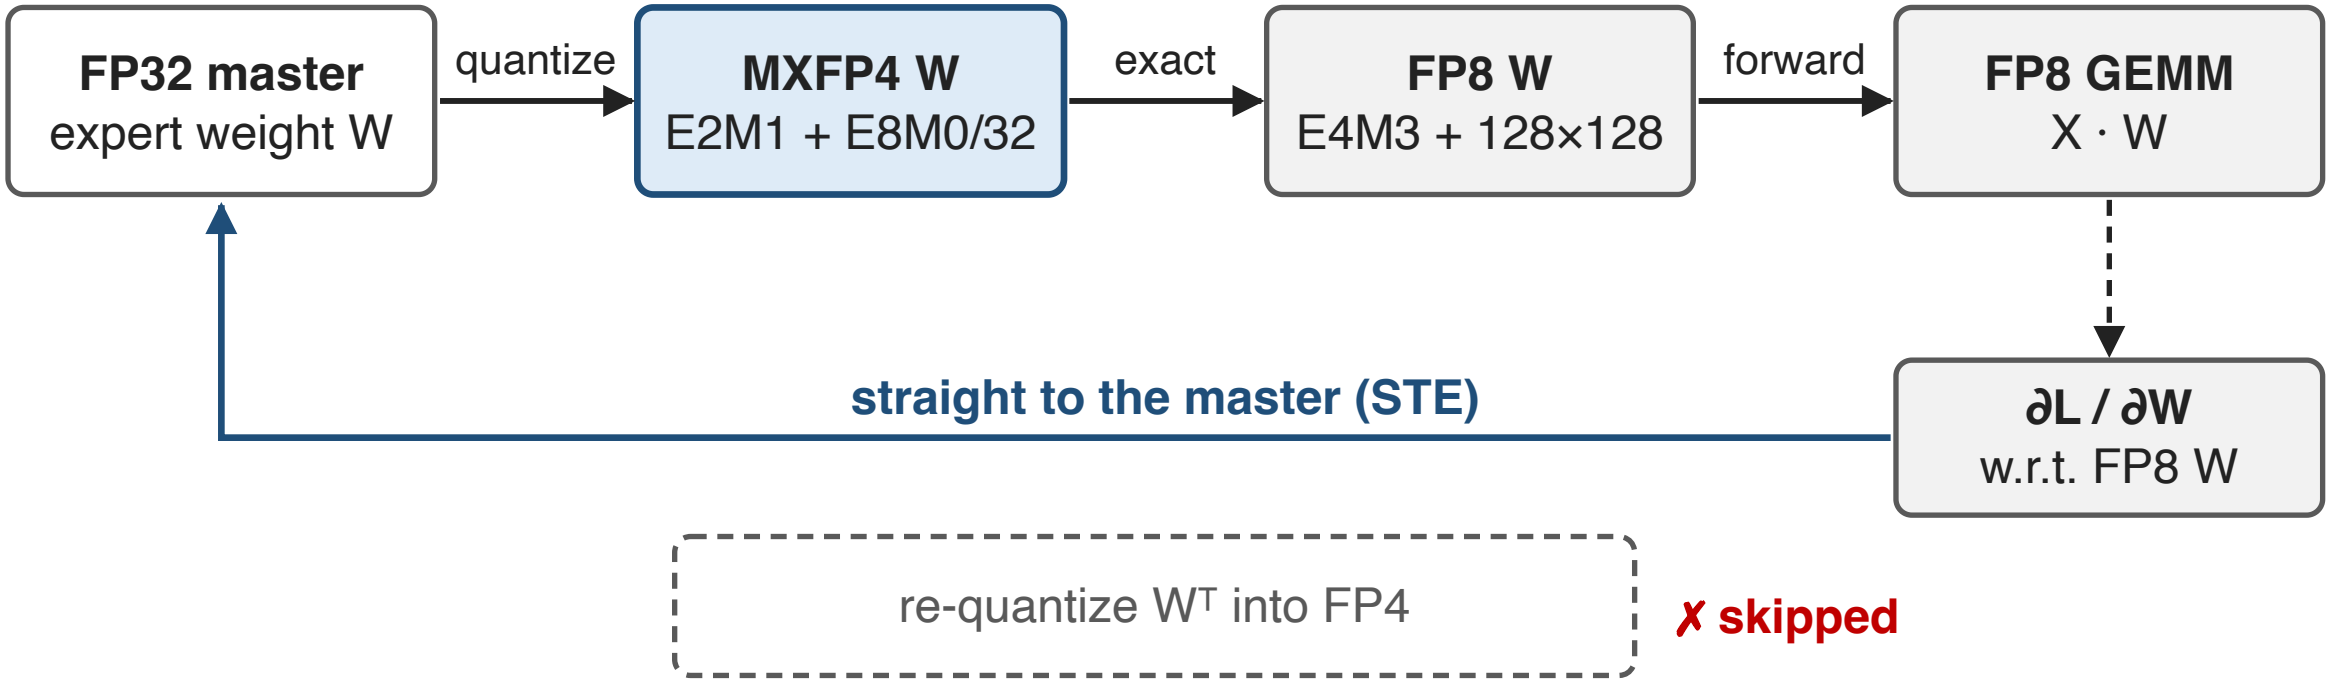
FP32 master
expert weight W
quantize
MXFP4 W
E2M1 + E8M0/32
exact
FP8 W
E4M3 + 128×128
forward
FP8 GEMM
X · W
straight to the master (STE)
∂L / ∂W
w.r.t. FP8 W
re-quantize W^T into FP4
✗ skipped