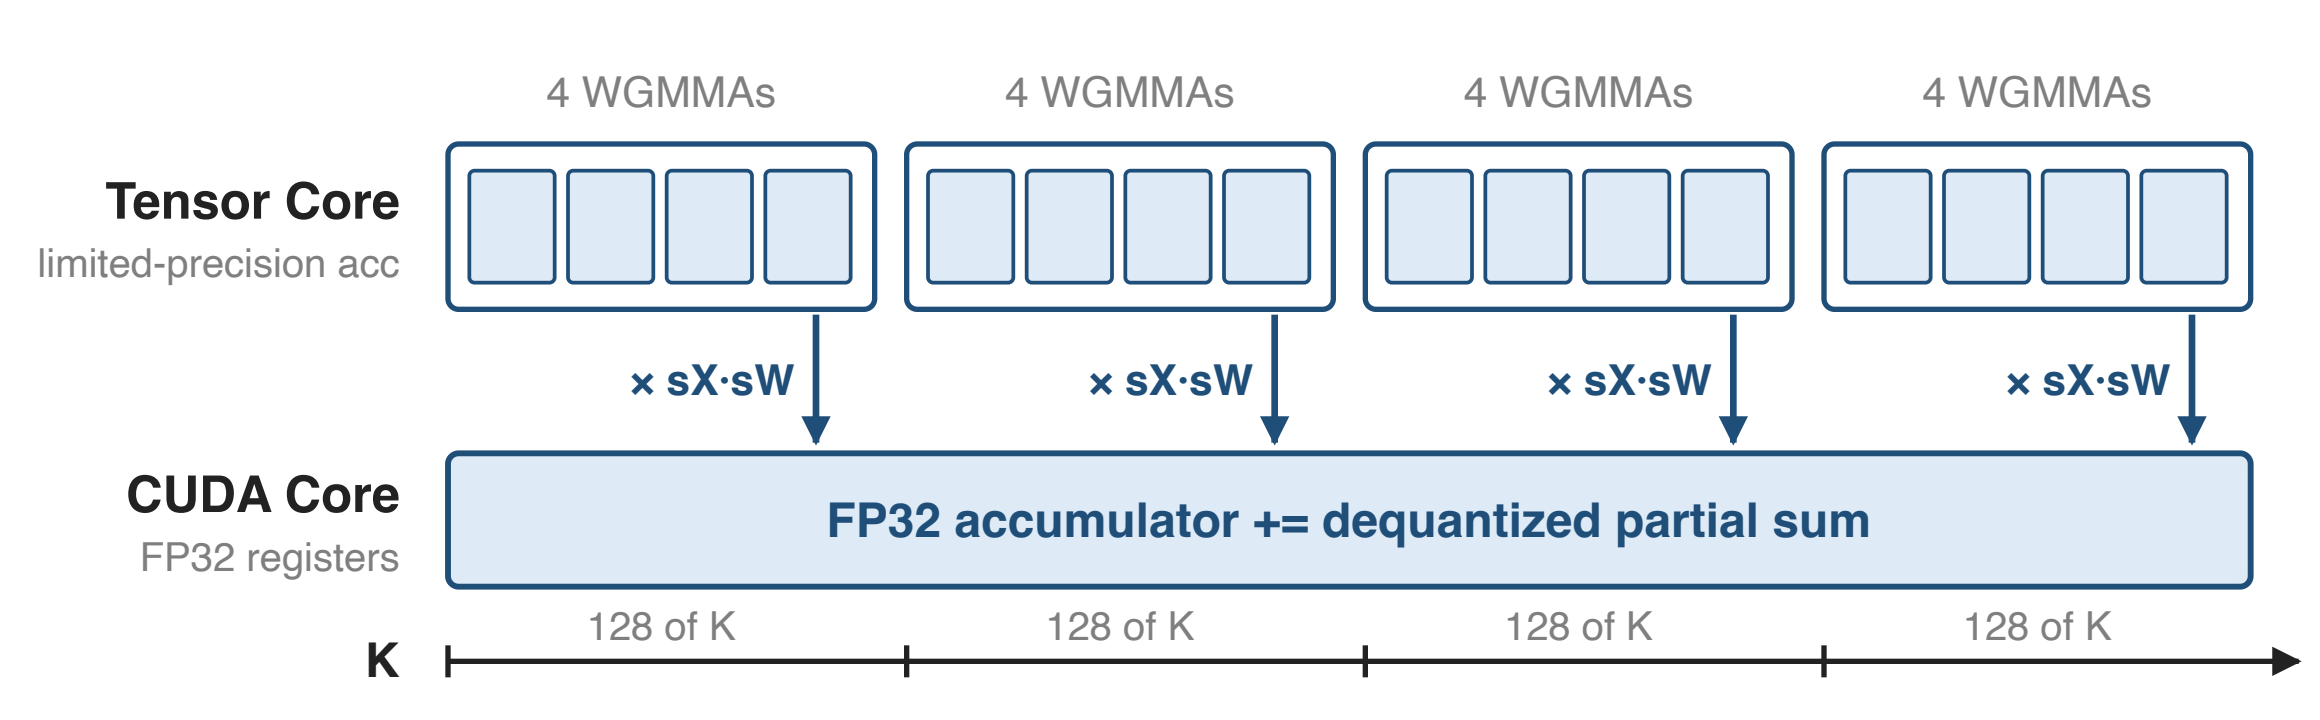
4 WGMMAs
4 WGMMAs
4 WGMMAs
4 WGMMAs
Tensor Core
limited-precision acc
× sX·sW
× sX·sW
× sX·sW
× sX·sW
CUDA Core
FP32 registers
FP32 accumulator += dequantized partial sum
128 of K
128 of K
128 of K
128 of K
K