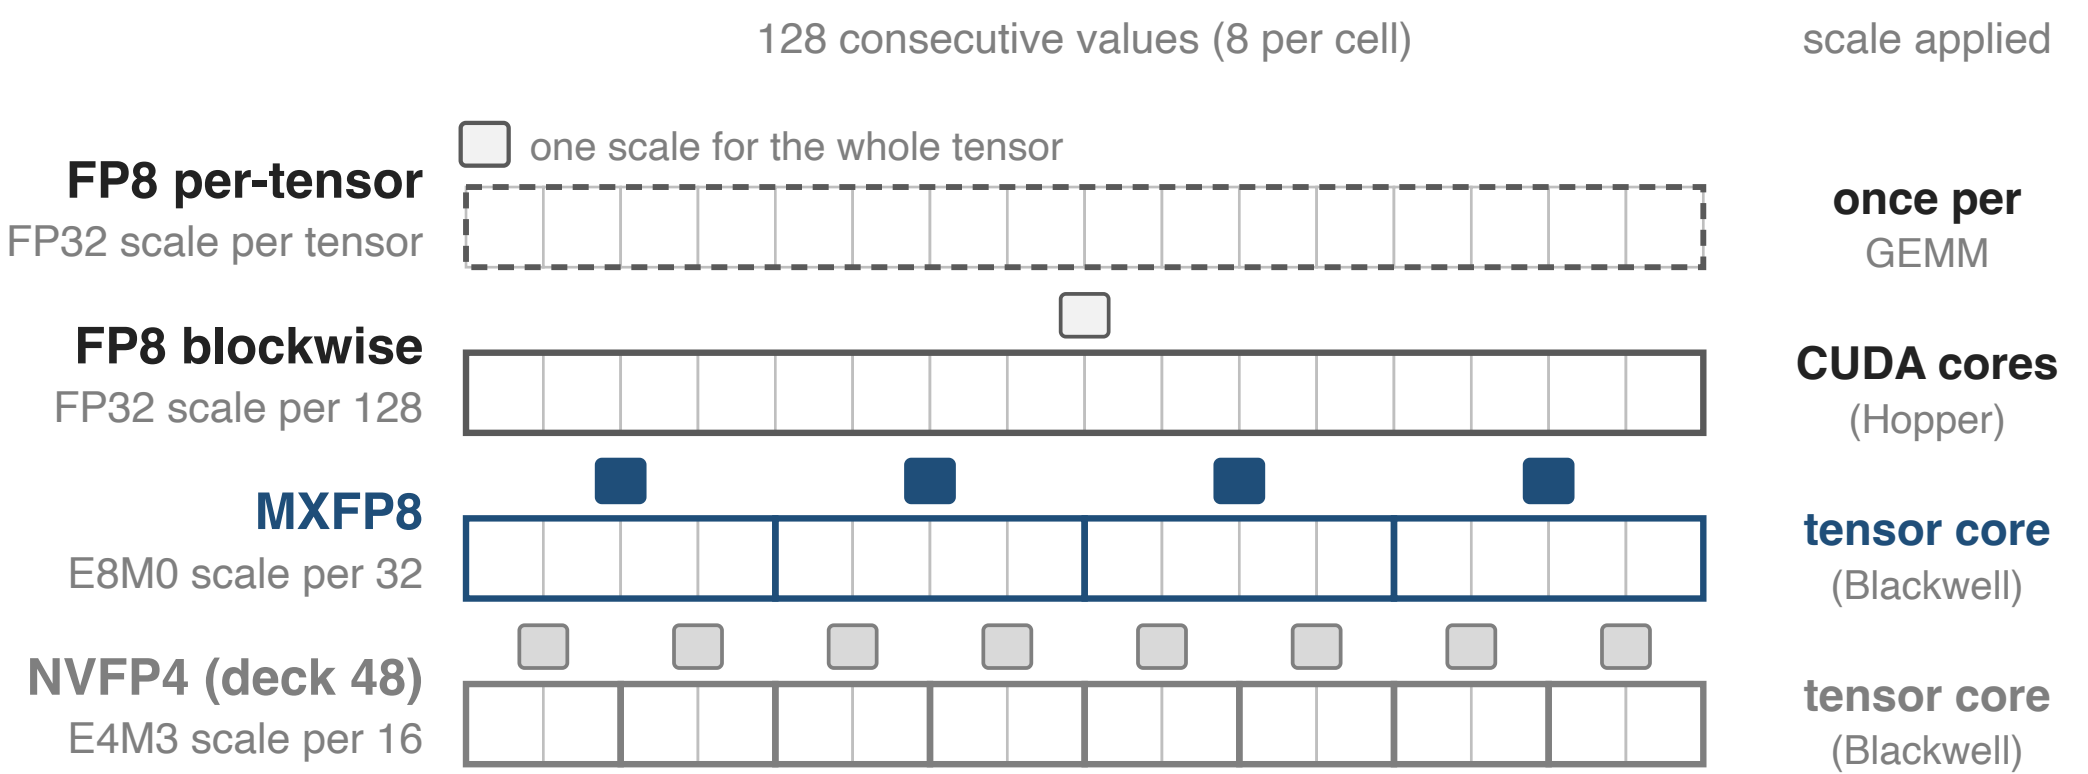
128 consecutive values (8 per cell)
scale applied
one scale for the whole tensor
FP8 per-tensor
FP32 scale per tensor
once per
GEMM
FP8 blockwise
FP32 scale per 128
CUDA cores
(Hopper)
MXFP8
E8M0 scale per 32
tensor core
(Blackwell)
NVFP4 (deck 48)
E4M3 scale per 16
tensor core
(Blackwell)