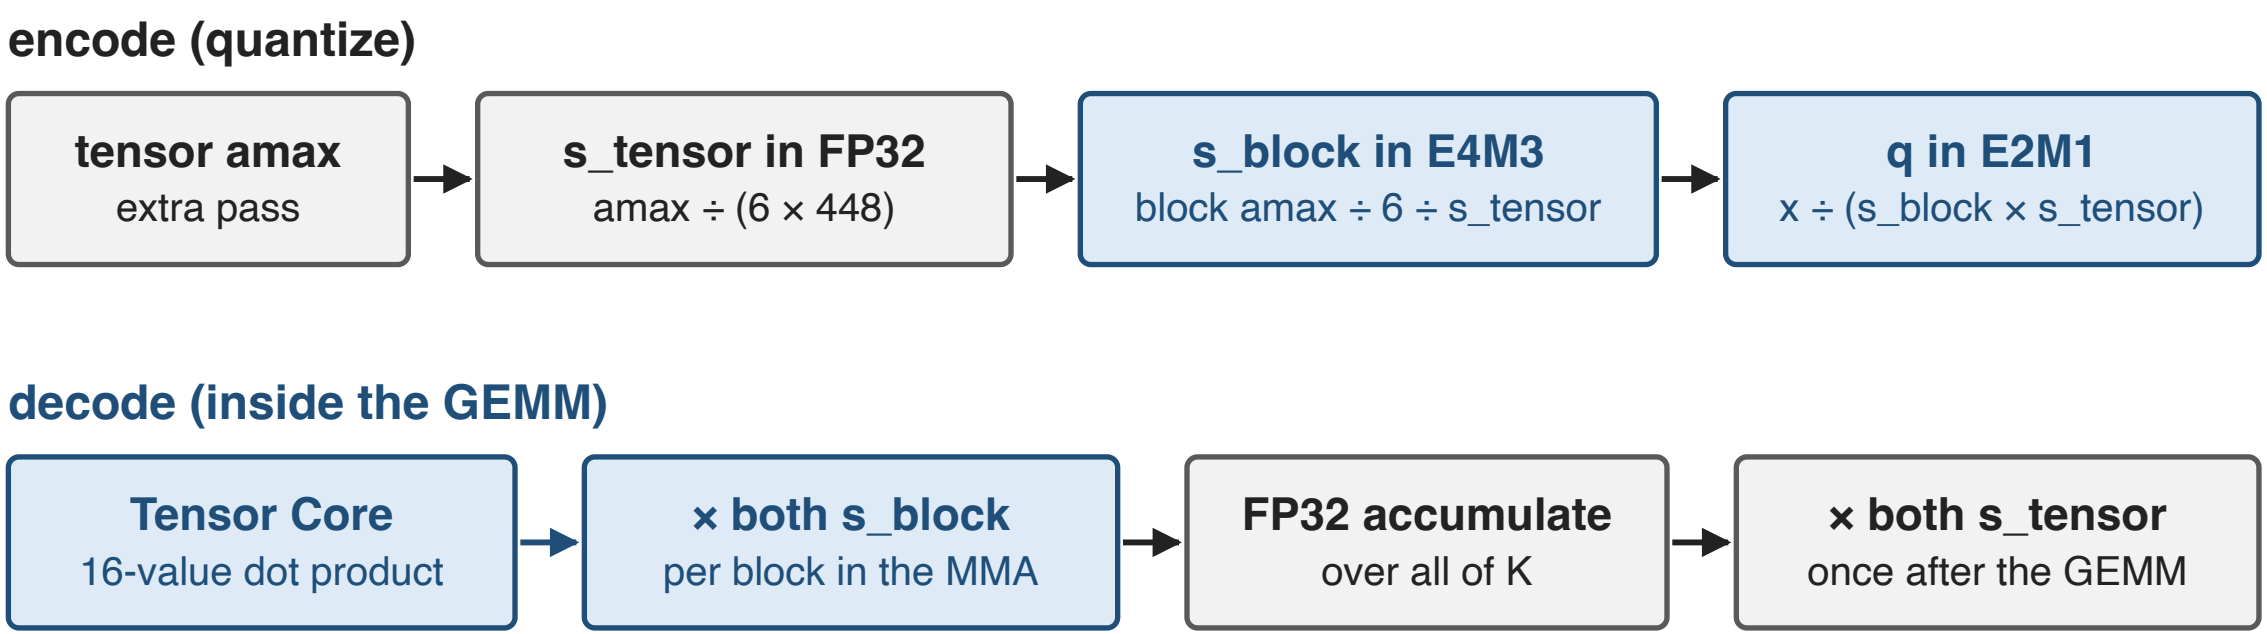
encode (quantize)
tensor amax
extra pass
s_tensor in FP32
amax ÷ (6 × 448)
s_block in E4M3
block amax ÷ 6 ÷ s_tensor
q in E2M1
x ÷ (s_block × s_tensor)
decode (inside the GEMM)
Tensor Core
16-value dot product
× both s_block
per block in the MMA
FP32 accumulate
over all of K
× both s_tensor
once after the GEMM