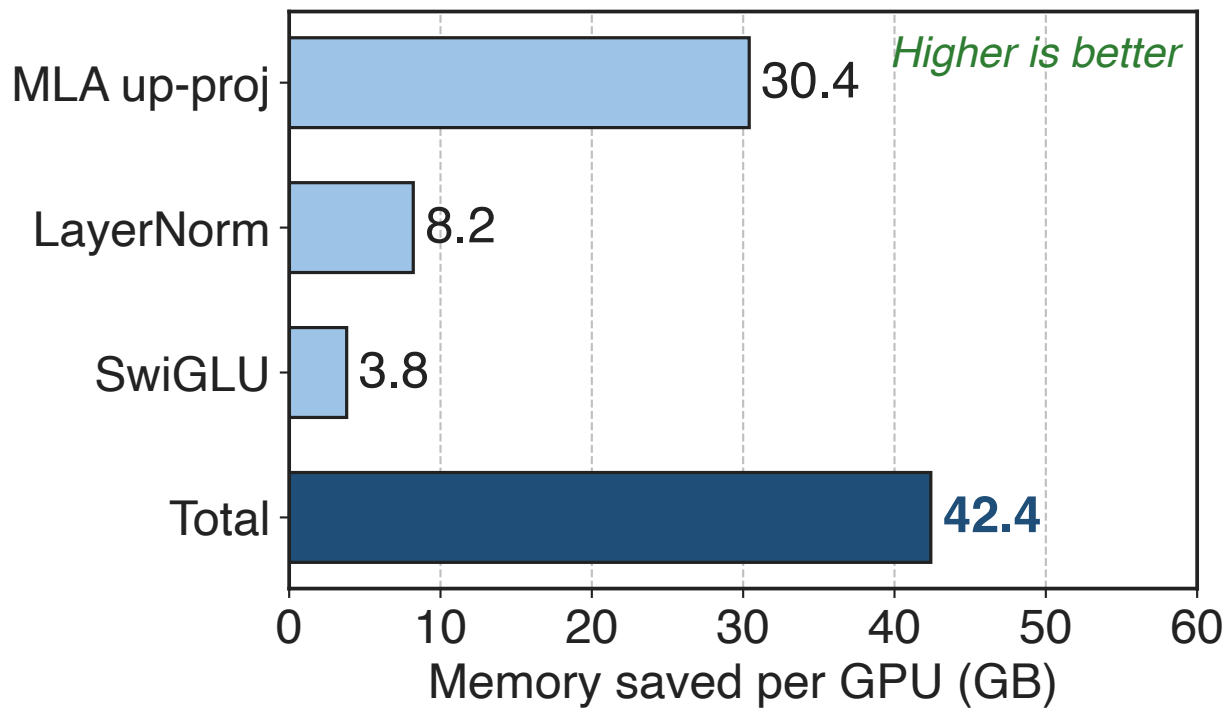
Higher is better
MLA up-proj
LayerNorm
SwiGLU
Total
30.4
8.2
3.8
42.4
0
10
20
30
40
50
60
Memory saved per GPU (GB)

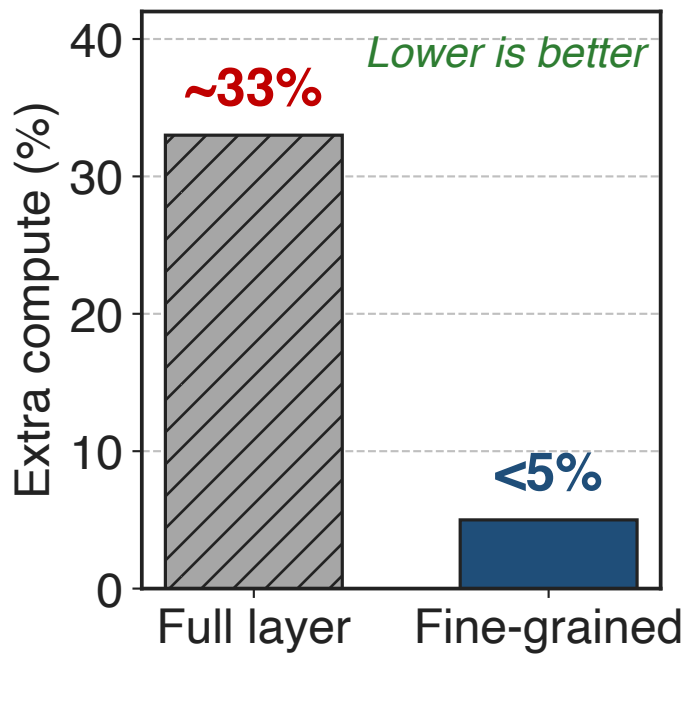
Lower is better
~33%
<5%
Extra compute (%)
0
10
20
30
40
Full layer
Fine-grained