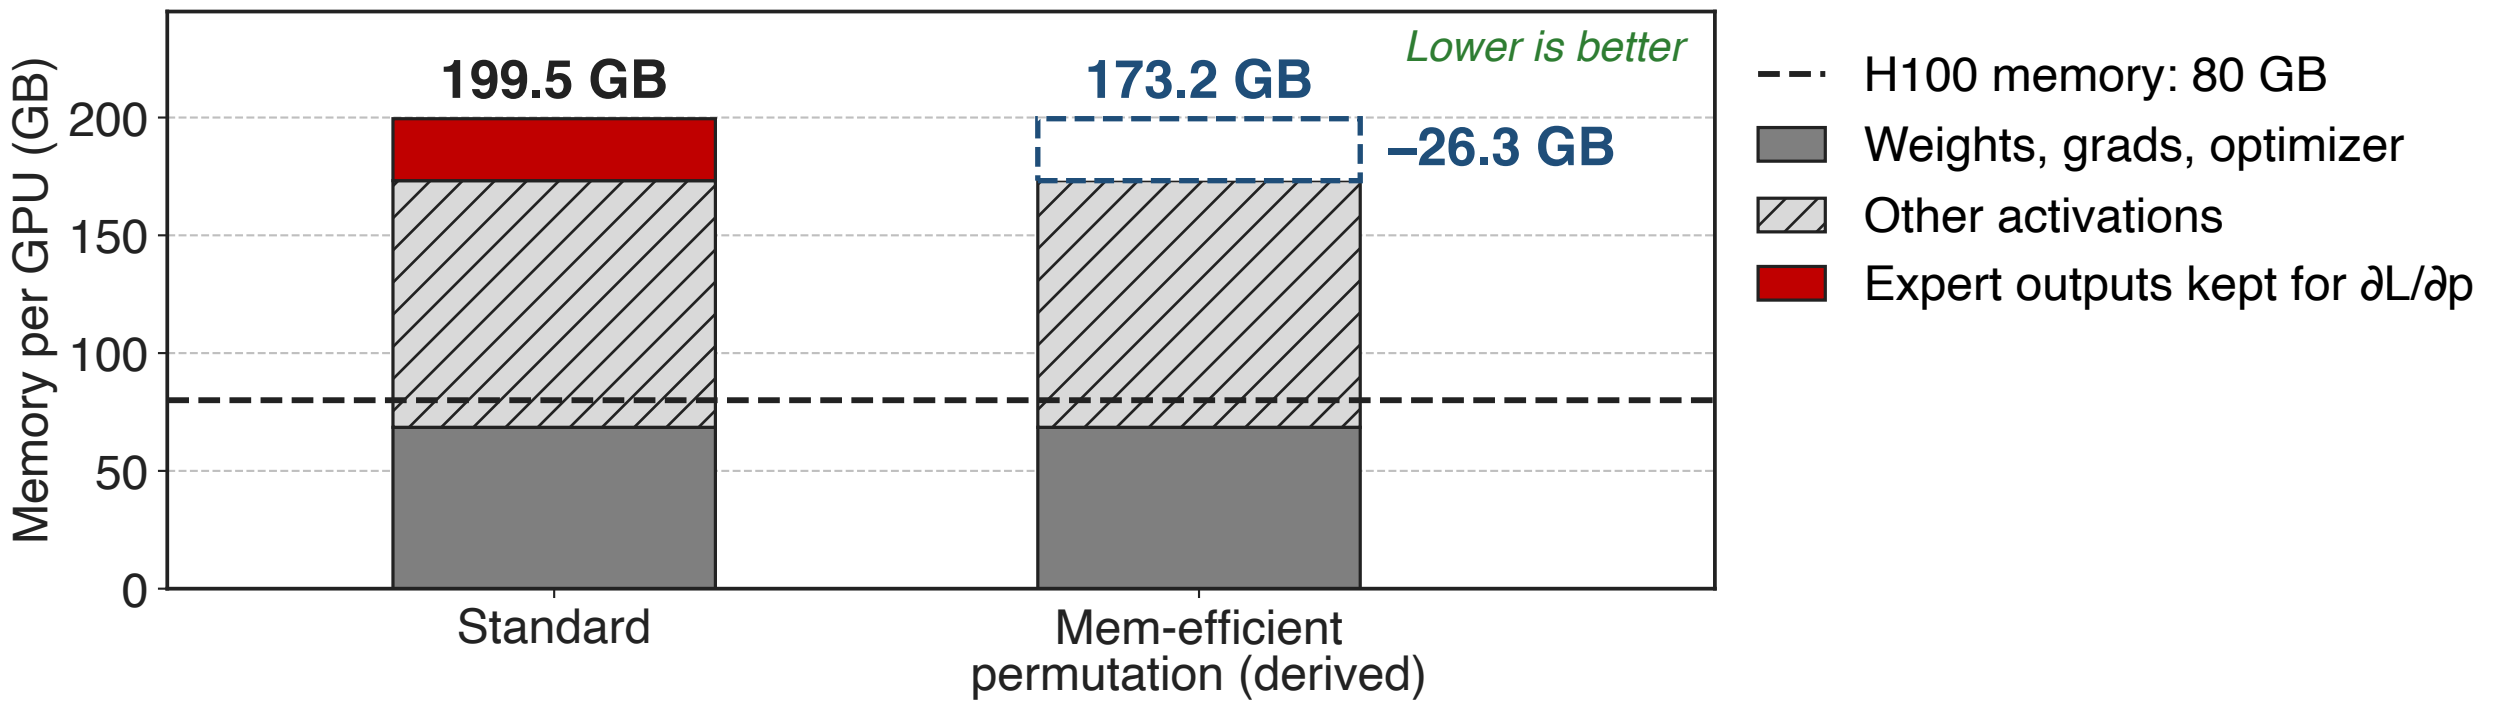
Lower is better
199.5 GB
173.2 GB
–26.3 GB
Memory per GPU (GB)
200
150
100
50
0
Standard
Mem-efficient
permutation (derived)
H100 memory: 80 GB
Weights, grads, optimizer
Other activations
Expert outputs kept for ∂L/∂p