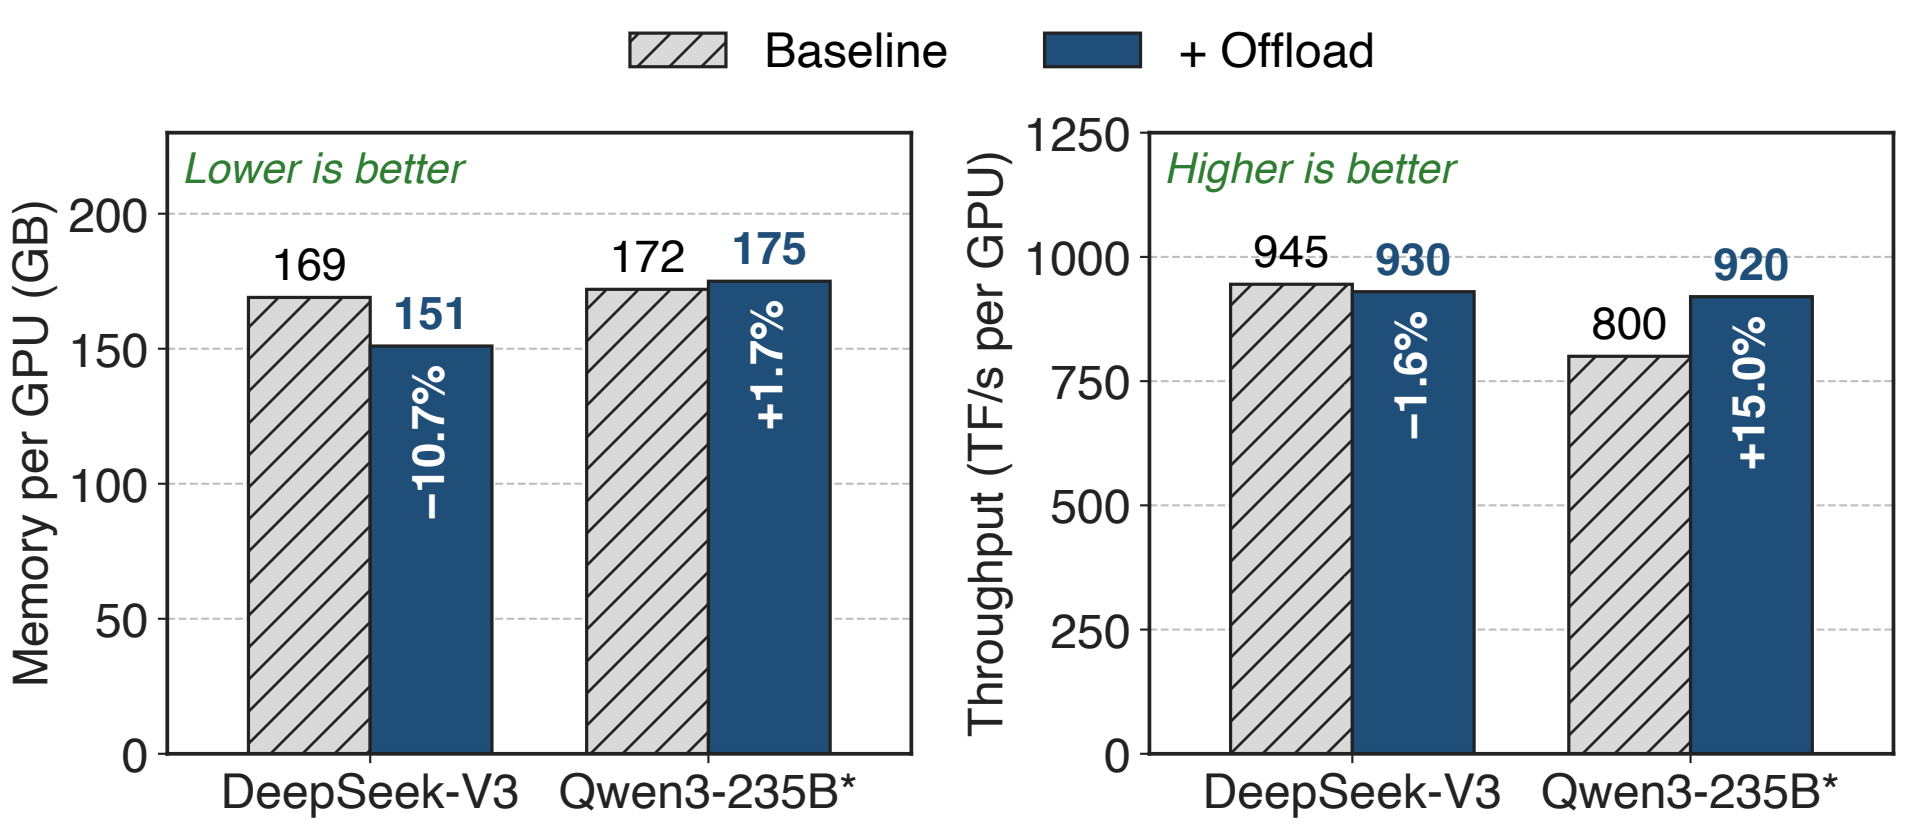
Baseline
+ Offload
Lower is better
Memory per GPU (GB)
200
150
100
50
0
169
151
−10.7%
172
175
+1.7%
DeepSeek-V3
Qwen3-235B*
Higher is better
Throughput (TF/s per GPU)
1250
1000
750
500
250
0
945
930
−1.6%
800
920
+15.0%
DeepSeek-V3
Qwen3-235B*