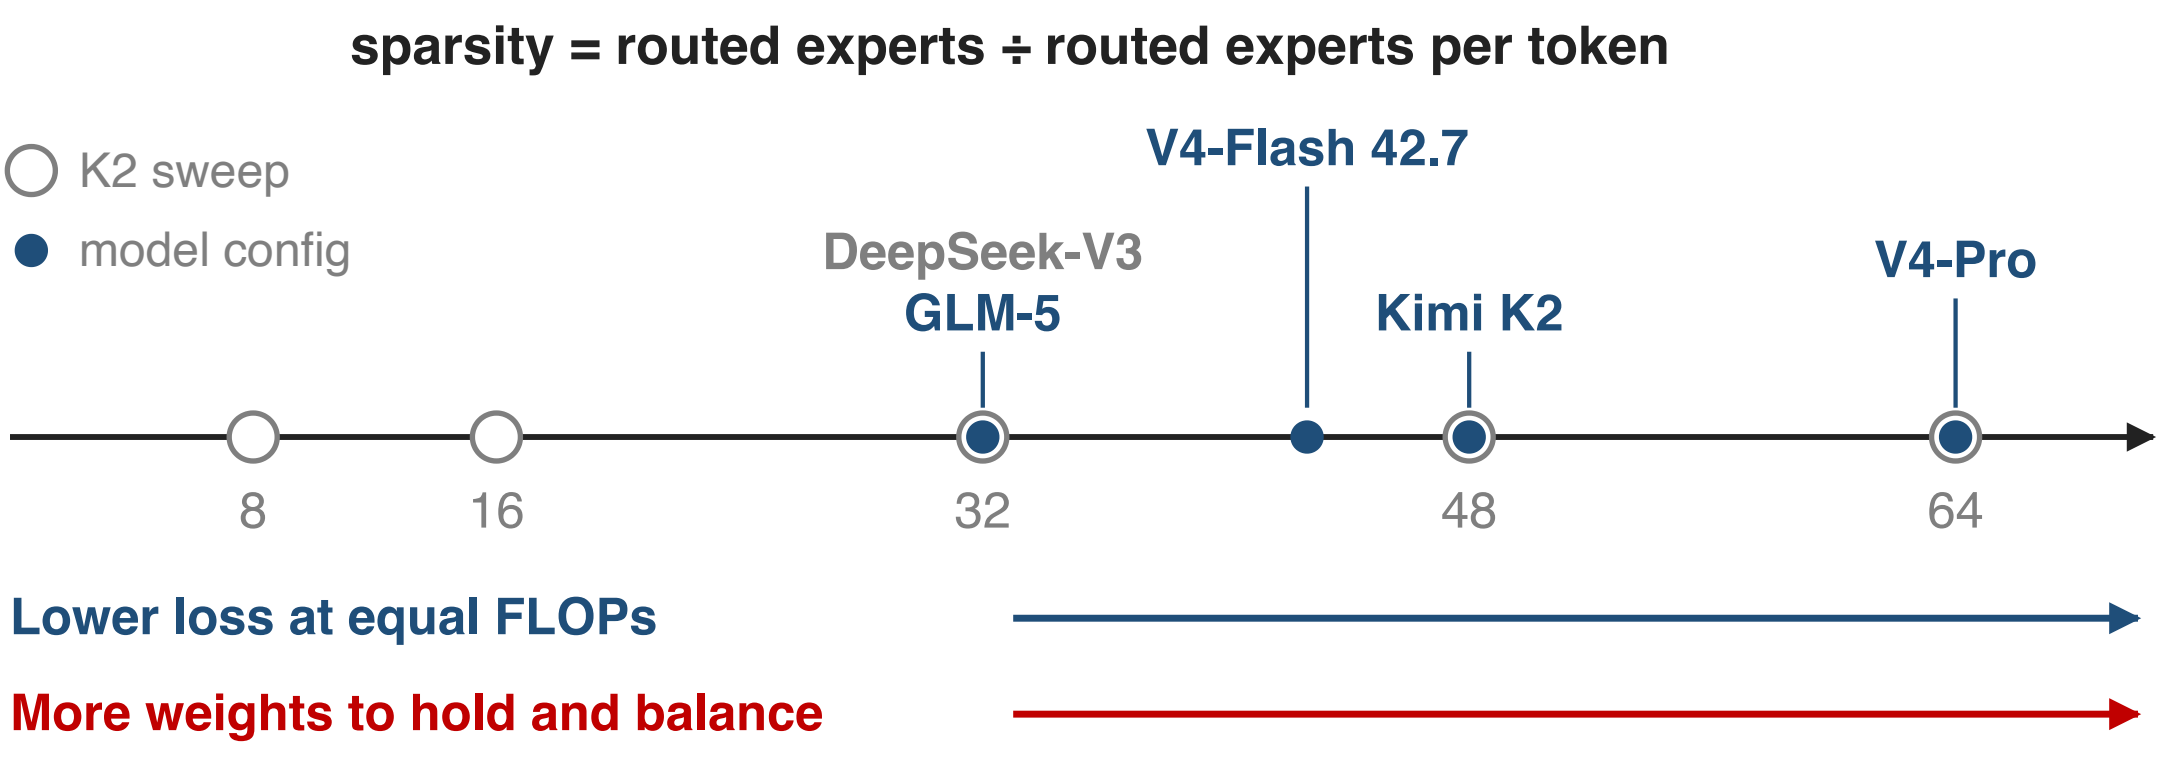
sparsity = routed experts ÷ routed experts per token
K2 sweep
model config
V4-Flash 42.7
DeepSeek-V3
GLM-5
Kimi K2
V4-Pro
8
16
32
48
64
Lower loss at equal FLOPs
More weights to hold and balance