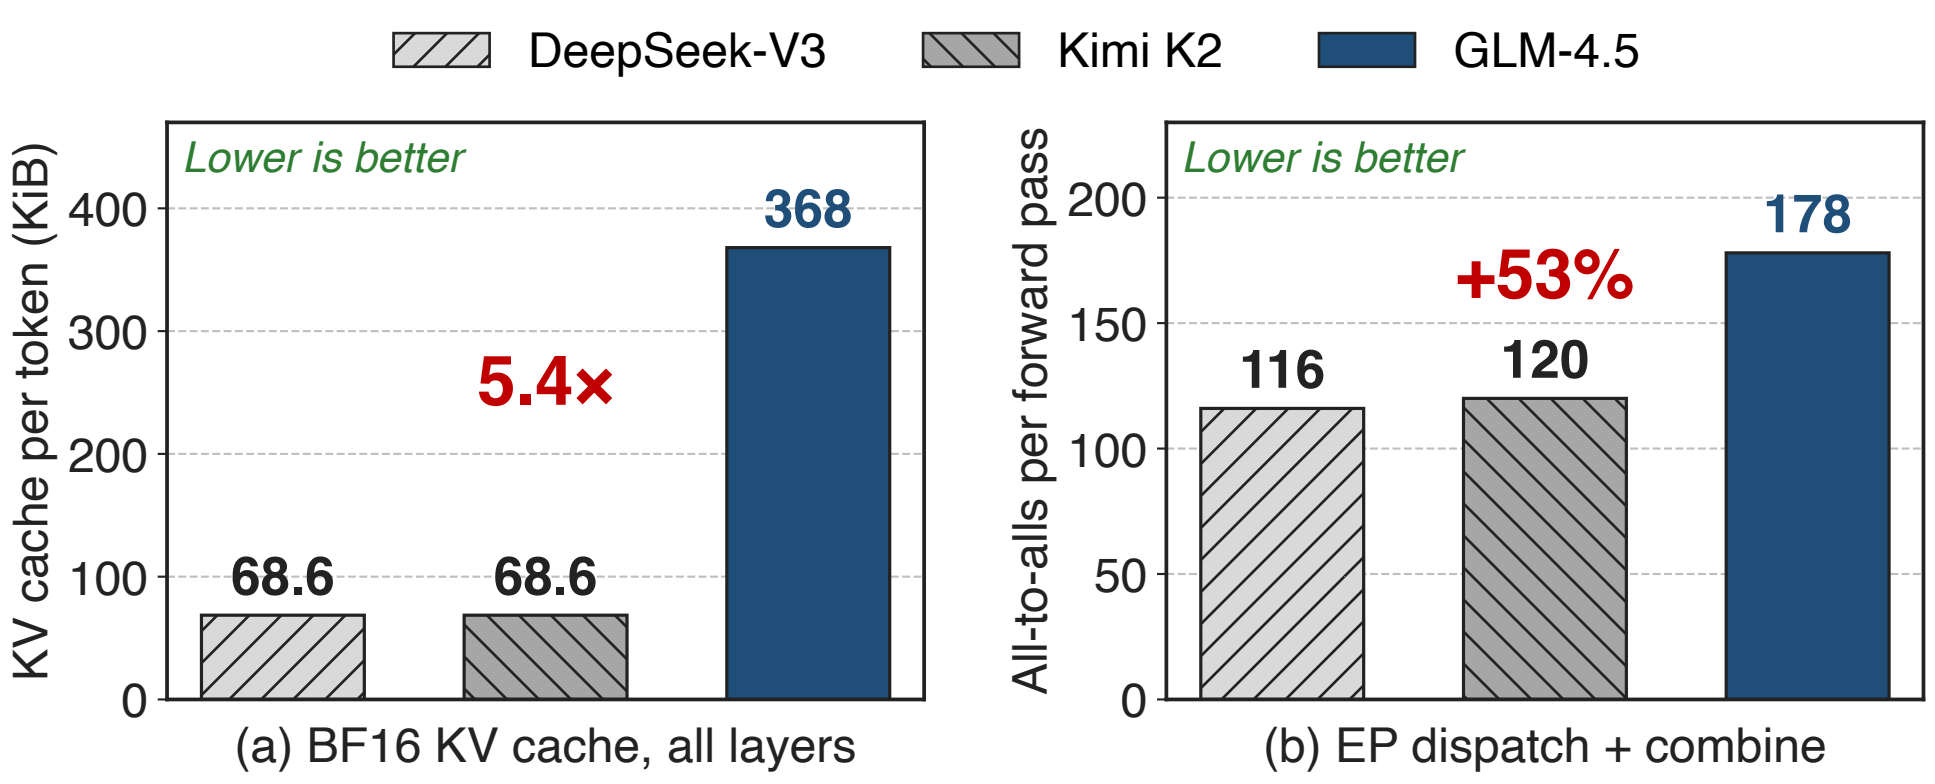
DeepSeek-V3
Kimi K2
GLM-4.5
Lower is better
KV cache per token (KiB)
400
300
200
100
0
368
5.4×
68.6
68.6
(a) BF16 KV cache, all layers
Lower is better
All-to-alls per forward pass
200
150
100
50
0
178
+53%
116
120
(b) EP dispatch + combine