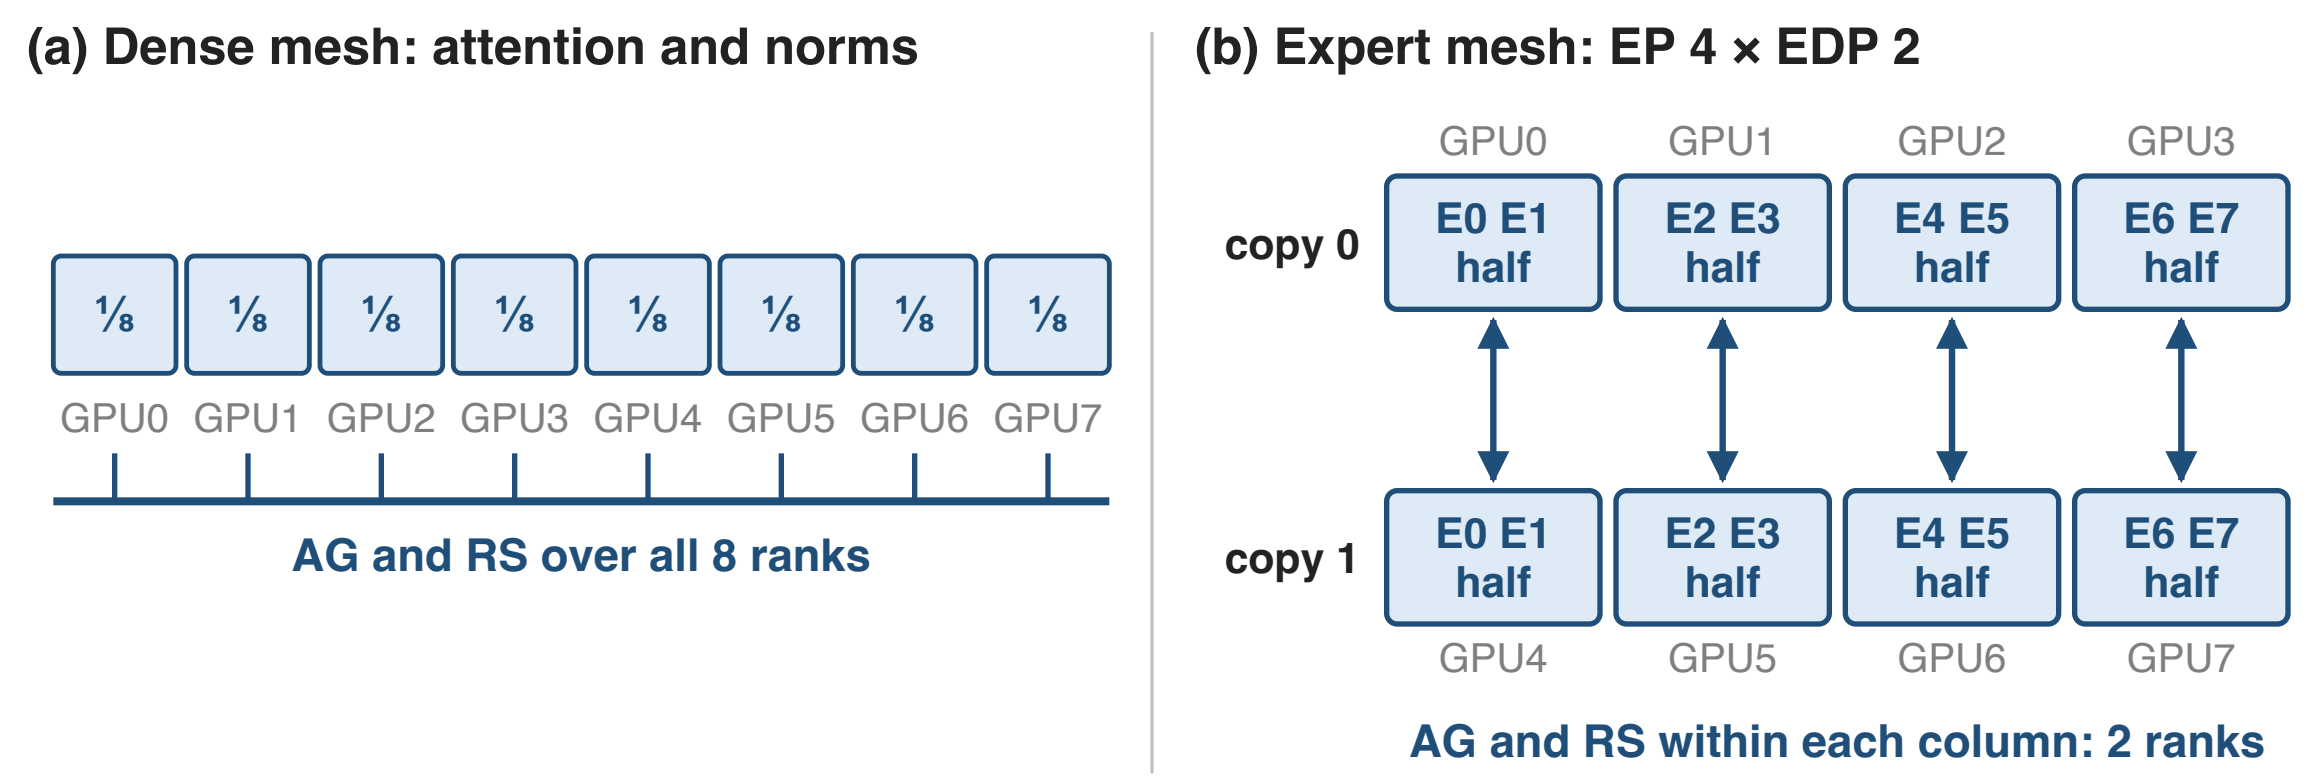
(a) Dense mesh: attention and norms
⅛
⅛
⅛
⅛
⅛
⅛
⅛
⅛
GPU0
GPU1
GPU2
GPU3
GPU4
GPU5
GPU6
GPU7
AG and RS over all 8 ranks
(b) Expert mesh: EP 4 × EDP 2
GPU0
GPU1
GPU2
GPU3
copy 0
E0 E1 half
E2 E3 half
E4 E5 half
E6 E7 half
copy 1
E0 E1 half
E2 E3 half
E4 E5 half
E6 E7 half
GPU4
GPU5
GPU6
GPU7
AG and RS within each column: 2 ranks