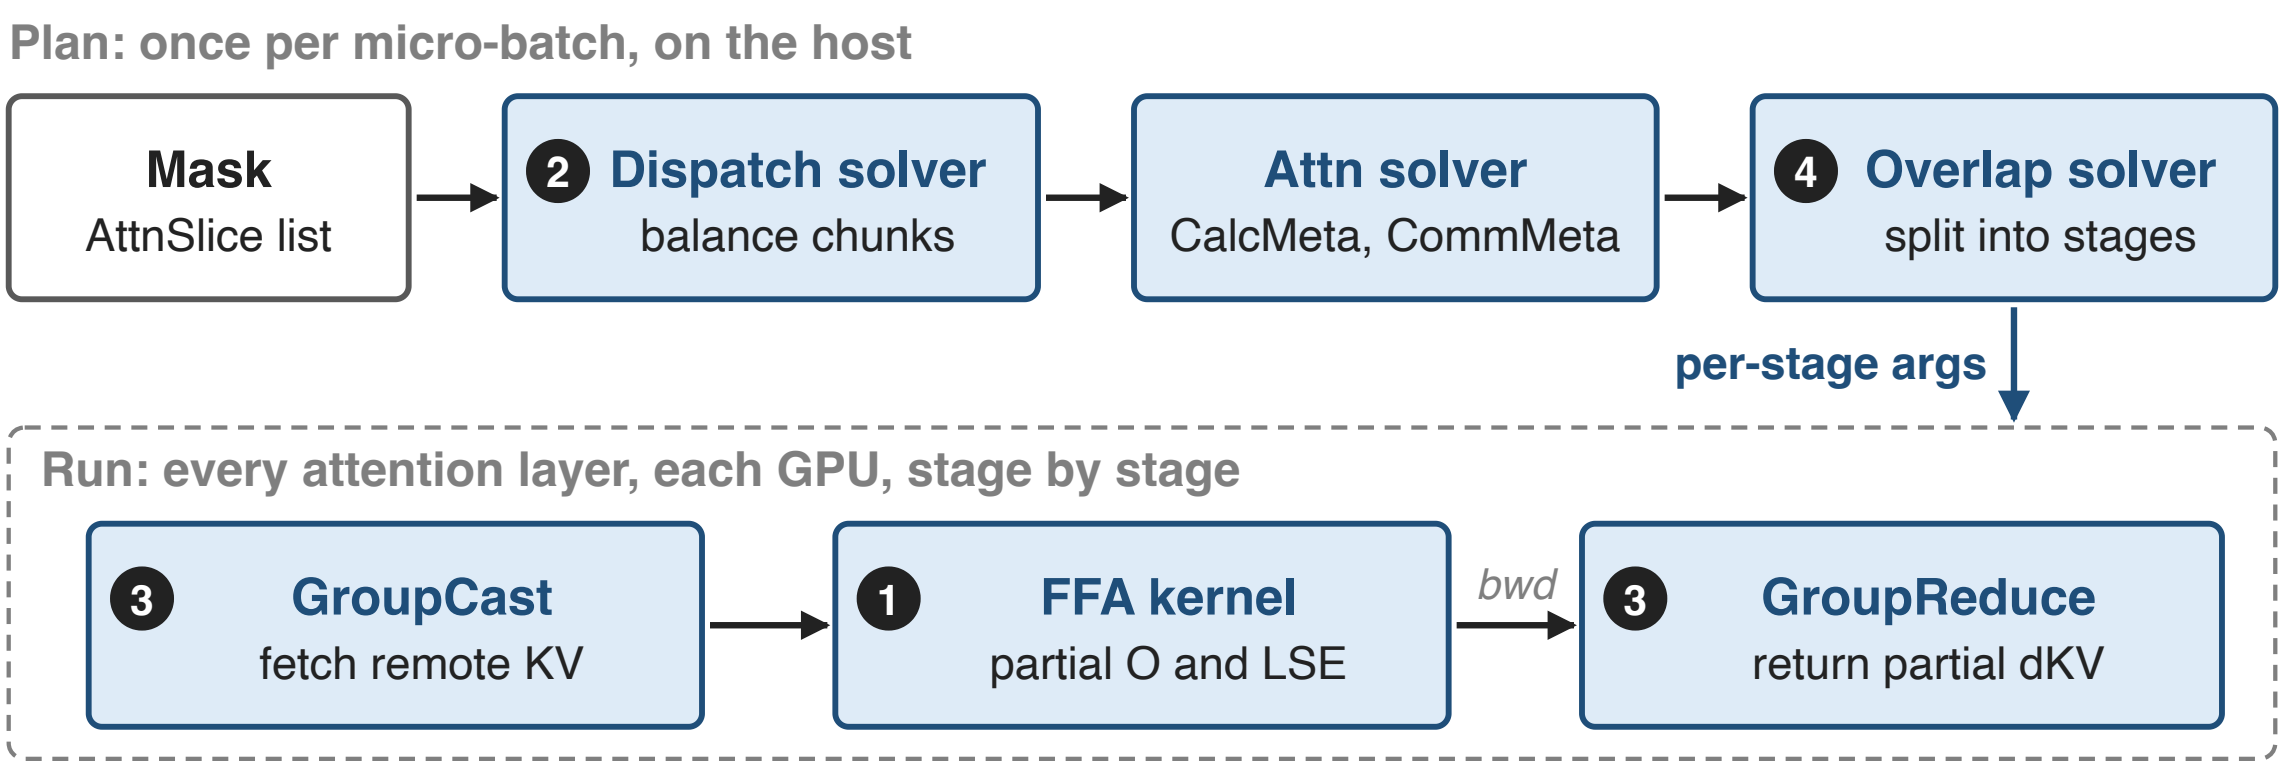
Plan: once per micro-batch, on the host
Mask
AttnSlice list
2
Dispatch solver
balance chunks
Attn solver
CalcMeta, CommMeta
4
Overlap solver
split into stages
per-stage args
Run: every attention layer, each GPU, stage by stage
3
GroupCast
fetch remote KV
1
FFA kernel
partial O and LSE
bwd
3
GroupReduce
return partial dKV